

Dashboard Regression + Host-Render Evidence — 2026-07-10

Revision: 1 Last modified: 2026-07-09T20:50:11Z

Stream: DASH (T1/main). Scope: dashboard/ only. Conductor performs commit/push; this document + the raw logs beside it are the captured evidence per §11.4.5 / §11.4.6 / §11.4.107 / §11.4.170.

Environment: Node v24.18.0, npm 11.16.0, Playwright 1.60.0
(Chromium installed at ~/.cache/ms-playwright). See 00_env.log.

Summary table

#	Check	Command run	Result (real numbers)	Evidence
1	Type-check	<code>npm run -s typecheck (= tsc --noEmit -p tsconfig.json && tsc --noEmit -p tsconfig.node.json && tsc --noEmit -p tsconfig.vitest.json)</code>	0 error TS occurrences; process exit code 0; no stdout/stderr emitted (tsc silent-success)	01_type
2	Unit/ component tests	<code>npm run -s test:run (= vitest run)</code>	12 test files passed (12), 107 tests passed (107), 0 failed; verbatim summary line Test Files 12 passed (12) / Tests 107 passed (107); Duration 10.12s; exit code 0	02_unit
3				03_host

#	Check	Command run	Result (real numbers)	Evidence
			117 passed (117), 0 failed, exit code 0. Covers 13 screen/state surfaces (login, overview, deployments, fleet, fleet- empty, fleet- error, groups, audit, releases, releases-empty, releases-error, artifact-upload, appshell) × {light, dark} = 26 screen×theme renders, each carrying 4 checks (golden-baseline match via toHaveScreenshot, a self-validated layout/OCR oracle that PASSES on the good render and FAIL-detects on a deliberately mutated render, a self-validated pixelmatch image-diff analyzer with the same good/ mutated discrimination, and a check that the committed golden baseline REJECTS a mutated render)	
	Host-rendered pixel proof (\$11.4.170)	npx playwright test -- config=playwright.hostrender.config.ts		

#	Check	Command run	Result (real numbers)	Evidence
			+ 1 light-vs-dark distinct-surface check per surface. Runtime 45.1s.	
4	Lint	none — no lint script in dashboard/ package.json, no ESLint config file present in dashboard/	N/A — verified absence via <code>grep -i '"lint"'</code> package.json (no match) + directory listing (no .eslintrc*)	00_env. presenc no separ needed. confirm

Notes

- All commands were run for real, from dashboard/ with the repo's own installed node_modules (no invented commands; every script above is copied verbatim from dashboard/package.json).
- No regression found. Every real check that exists in this project (typecheck, unit/component tests, host-render visual-proof suite) is green with real captured numbers, not metadata-only claims.
- The host-render suite is genuinely self-validating: each screen's "layout oracle" and "image-diff analyzer" checks assert PASS on the known-good render AND FAIL-detection on a deliberately mutated render in the same test run — this is the §11.4.107(10) / §11.4.170 self-validated golden-good/golden-bad discipline, not a bare screenshot diff.
- Playwright's webServer (vite on port 4318, reuseExistingServer: false) was started and torn down by the Playwright test runner itself; a post-run process check (ps aux | grep 'vite --port 4318') found no orphaned process — target left quiescent per §11.4.14.
- No git command was run by this stream (conductor-owned per instructions).
- HTML/PDF siblings for this document are intentionally NOT generated here — the conductor exports them at commit time per §11.4.65/§11.4.106.

Raw evidence files in this directory

- 00_env.log — tool-version + tool-presence capture (node, npm, tsc, vitest, playwright binaries)
- 01_typecheck.log — full npm run -s typecheck output + exit code
- 02_unit_tests.log — full npm run -s test:run output + exit code
- 03_hostrender.log — full npx playwright test --config=playwright.hostrender.config.ts output + exit code