

Target Multi-Account (Tenant) Data Model — To-Be Design

Revision: 1 **Last modified:** 2026-07-10T11:18:54Z

Scope. This is the KEYSTONE data-model design for adding an ACCOUNT (tenant) layer to the Helix OTA control plane (server/, Go/Gin modular monolith). It designs the to-be entities (Account, User, Membership, Role/Permission), how they scope the existing OTA model (Device/Artifact/Release/Deployment/Telemetry/Group), the store.Repository extension for BOTH the in-memory MVP and the pgx target, a migration-style DDL sketch, the OQ-4 migration path, and the super-admin at the data layer. It is the single source of truth the sibling to-be docs (21_authz_rbac_superadmin, 22_api_surface, 24_project_side_cli, 26_security_threat_model, 30_delivery_plan) reference — so it is deliberately precise and self-consistent (§11.4.186 anti-divergence).

Reading order. Read 00_INDEX.md (mandate + open questions) and 10_existing_auth_and_project_model.md (the authoritative as-is facts) FIRST. Every as-is claim below is grounded in a file:line those docs already established; this doc does not re-derive them, it cites them. Every design choice is grounded in a cited source (## Sources verified 2026-07-10) or marked “original work”. Every open choice is presented as a recommendation-with-tradeoffs for operator decision — never a silent decision (§11.4.6).

0. The one decision this doc anchors: tenancy model

Three canonical multi-tenant storage models exist (Microsoft Azure Architecture Center; AWS; dasroot 2026): **(1) shared database + shared schema** with a `tenant_id/account_id` column on every tenant-owned row and database-enforced row scoping; **(2) schema-per-tenant**; **(3) database-per-tenant**.

Model	Isolation	Ops cost	Cross-tenant analytics	Fit for Helix OTA today
Shared schema + <code>account_id</code> (+ RLS)	Policy-enforced; a policy/where bug is a leak	Lowest — one schema, one migration	Trivial (one table)	Best — matches ADR-0003 “modular monolith — one PostgreSQL schema” (<code>001_initial_schema.up.sql:16</code>) and the existing single-schema store (<code>schema_postgres.sql:11</code>)
Schema-per-tenant	Strong; per-tenant customization	Every migration ×N tenants	Harder	Overkill; no per-tenant schema customization is required by the mandate
Database-per-tenant	Perfect; per-tenant residency	N connection pools, N backups	Hardest	Reserve for a future compliance/residency tenant only

Recommendation (KEYSTONE): shared schema + a denormalized `account_id` column on every tenant-owned table, defended by PostgreSQL Row-Level Security (RLS) at the `pgx` layer. This is the industry default when tenant counts can grow and no per-tenant schema customization is required (Microsoft, AWS, dasroot all recommend shared-schema+RLS as the default; escalate only for compliance/residency). It also requires the *least* change to the shipped single-schema store.

Tradeoff to accept: isolation now depends on correct scoping + correct RLS policies; a forgotten `WHERE account_id = ...` in application code would leak cross-tenant — which is exactly why RLS (database-enforced, “works even when your application code has bugs” — AWS) is the belt-and-suspenders second layer, and why `26_security_threat_model.md` must carry a paired mutation proving a non-scoped session cannot read another account’s rows.

Escalation path (operator decision, not now): a single high-compliance tenant that needs data residency or a separate backup schedule can later be promoted to its own schema or database without changing the shared-schema code for everyone else — the `account_id` seam is forward-compatible with a routing layer that maps some accounts to isolated storage.

1. Entities

Four new/lifted entities. Id type is **opaque TEXT** to match the *shipped* store convention (`project_id`, `device_id`, `artifact_id` are all TEXT — `schema_postgres.sql:15,34,49,240`), NOT the UUID used by the canonical design-target migration (`001_initial_schema.up.sql:40` etc.).

Divergence to reconcile (§11.4.186). The shipped store (`schema_postgres.sql`) uses opaque TEXT ids; the canonical target migration (`001_initial_schema.up.sql`) uses UUID and already declares `users + api_keys` tables (`10_* §3`, `001_initial_schema.up.sql:39-71`) that were never carried into the shipped store. **Recommendation:** the *implementation* seam (this design, §3–§4) uses opaque TEXT to mirror the store it actually extends; the canonical 001-style schema keeps UUID. Both remain internally consistent; §4 states the mapping so the two representations never silently diverge. Tradeoff: TEXT ids forgo `gen_random_uuid()` server-side generation (ids are minted app-side, as they already are for projects/devices) — accepted, because it matches the existing code and avoids a store-wide id-type migration.

1.1 Account — the tenant (top of the tree)

The account is the top of the hierarchy `account → projects → OTA updates`. It has no parent. Account-level identifiers are legitimately **global-unique** (accounts ARE the tenant boundary, so there is no enclosing scope).

Field	Type	Notes
<code>account_id</code>	TEXT PK	opaque, app-minted (mirrors <code>project_id</code>)
<code>name</code>	TEXT NOT NULL	display name; UNIQUE (global) — one tenant per name in the super-admin console

Field	Type	Notes
slug	TEXT NOT NULL	URL-/token-safe stable handle; UNIQUE (global) ; the stable key an account switcher / CLI selects by (\$11.4.111 resolve-by-stable-name)
status	TEXT NOT NULL DEFAULT 'active'	CHECK IN ('active', 'suspended', 'archived') — suspend gates all sign-in/scoped access without deleting data
created_at	TIMESTAMPTZ NOT NULL	
updated_at	TIMESTAMPTZ NOT NULL	

- **PK:** account_id. **Unique:** name, slug (both global).
- Go type mirrors store.Project (store.go:289-295) one level up.

1.2 User — persisted identity (closes the \$3 gap)

Today there is **no persisted user** — identity is the bare token Subject string and an in-memory env-seeded map (10_* \$3; token.go:32, users.go:8-48, main.go:96-104). This design persists the User the target migration already sketched (001_initial_schema.up.sql:39-51). A User is a **single global identity** that can belong to many accounts (the mandate: “users belong to one OR more accounts”, 00_INDEX \$1.3) — so user identifiers are global-unique and the *account relationship* lives in Membership (\$1.3), not on the user row.

Field	Type	Notes
user_id	TEXT PK	opaque, app-minted
username	TEXT NOT NULL	login identity; UNIQUE (global) ; the value the token Subject resolves to
email	TEXT NOT NULL	UNIQUE (global)
password_hash	TEXT NOT NULL	credential hash only, never cleartext (mirrors target users.password_hash, 001_initial_schema.up.sql:43)
is_super_admin	BOOLEAN NOT NULL DEFAULT FALSE	the global super-admin flag (\$6) — bypasses account scoping; set ONLY via config/

Field	Type	Notes
		env bootstrap, never a request (§11.4.10; 10_* §7)
is_active	BOOLEAN NOT NULL DEFAULT TRUE	disable without delete
created_at / updated_at	TIMESTAMPTZ NOT NULL	

- **PK:** user_id. **Unique:** username, email.
- **Resolves the CallerID ambiguity** (10_* §8 assumption 5): the token Subject (a username) resolves to user_id; memberships (§1.3) then attach roles per account. ProjectAccess.CallerID (store.go:307-311) should migrate to reference user_id (reconciliation item, §2).

1.3 Membership — user ↔ account (the many-to-many)

A user belongs to one OR more accounts (00_INDEX §1.3); an account has many users. This is the direct generalization of the existing ProjectAccess (caller ↔ project ↔ role, store.go:306-311) lifted one level to account ↔ user ↔ role. It is the table the WorkOS multi-tenant RBAC model calls the membership join (user_roles keyed (user_id, tenant_id, role_id)).

Field	Type	Notes
user_id	TEXT NOT NULL	FK → users(user_id) ON DELETE CASCADE
account_id	TEXT NOT NULL	FK → accounts(account_id) ON DELETE CASCADE
role	TEXT NOT NULL DEFAULT 'viewer'	per-account role (§1.4); CHECK IN ('viewer', 'operator', 'admin')
is_owner	BOOLEAN NOT NULL DEFAULT FALSE	the account's owner (the membership minted at bootstrap, §3.3); at most one owner per account is a policy 21_* decides
granted_at	TIMESTAMPTZ NOT NULL	
granted_by		

Field	Type	Notes
	TEXT NOT NULL DEFAULT ''	granting user_id (super-admin or account admin)

- **PK:** composite (user_id, account_id) — one membership row per (user, account); a user in N accounts has N rows. **Role is per-account**, so the same user can be admin in account A and viewer in account B (WorkOS: “the same user may hold different roles in different contexts”).
- Go type: AccountMembership{AccountID, UserID string; Role AccountRole; IsOwner bool}.

1.4 Role / Permission — the per-account authorization vocabulary

The mandate asks for “ultimate flexibility to the maximum” (00_INDEX §1.6), but the *shape* of the permission model (RBAC vs ABAC vs role+scope hybrid) is **OQ-2, explicitly owned by 21_authz_rbac_superuseradmin.md** (00_INDEX §5). This keystone doc models only the minimal role skeleton the data layer needs and reserves the richer path so 21_* can land it without a second data-model change.

Recommendation for THIS doc (defer full shape to 21_*): ship the fixed per-account role enum on Membership.role now, mirroring the existing ProjectRole hierarchy viewer < operator < admin (store.go:297-304) so the existing rank-compare tooling (handlers_project.go:78-92) reuses unchanged. AccountRole = viewer < operator < admin, distinct from but structurally identical to ProjectRole.

Reserved “max-flexibility” path (roles/permissions tables), if 21_* chooses it: a tenant-scoped RBAC set — roles(role_id, account_id, name, UNIQUE(account_id,name)), permissions(permission_id, action, resource), role_permissions(role_id, permission_id), and Membership.role becoming a role_id FK — the “hybrid/templates” model WorkOS recommends (default roles + per-tenant customization). **Tradeoff:** fixed-enum is simplest and reuses existing tooling but caps flexibility at three roles; the roles/permissions tables give the mandated “ultimate flexibility” at the cost of a role-resolution join on every authZ check. Presented for 21_*; §4 emits the fixed-enum column now + the reserved tables as commented DDL so either path is a pure additive migration.

1.5 How Account sits ABOVE Project

Account (tenant)	accounts
└─ Project (add account_id FK)	projects.account_id →
accounts	
└─ OTA updates:	(each carries account_id
[RLS key] + project_id)	
Artifact, Release,	
Deployment, Device,	
Group, Telemetry	

The as-is gap (10_* §4): projects are **empty containers** — Device/Artifact/Release/Deployment/Telemetry/Group carry **no ProjectID** (store.go:36-178), so the account → projects → OTA updates link is entirely absent below the project row. §2 adds it.

2. Scoping the existing OTA model

Design rule (from §0 recommendation + AWS RLS best practice): every tenant-owned table gets a **denormalized account_id** (NOT NULL, the RLS scoping key, and the leading column of its primary access index — AWS: “every table with an RLS policy needs tenant_id as the leading column in its primary access indexes”), AND a project_id (the true mid-tier parent). account_id is denormalized onto each resource rather than reached only via project_id → account_id because RLS policies and the hot scoping filter must not pay a join on every row.

Tradeoff of the denormalization: resource.account_id must always equal project.account_id. Enforce with a **composite FK** (project_id, account_id) REFERENCES projects(project_id, account_id) (requires a UNIQUE(project_id, account_id) on projects) so the database — not application discipline — guarantees the two never drift. Alternative (a trigger) is weaker; the composite FK is “original work” applied to this schema and is the recommended enforcement.

2.1 Per-entity column plan

Entity (struct / table)	add account_id	add project_id	Nullability	Migration / backfill implication
	yes (it is the mid-	n/a	account_id NOT NULL	backfill every existing project to

Entity (struct / table)	add account_id	add project_id	Nullability	Migration / backfill implication
Project (store.go:289; projects schema_postgres.sql:239)	tier, not the top)		after backfill	the default account (5)
Device (store.go:36- *; devices :14)	yes	yes	account_id NOT NULL ; project_id NULL- until- assigned (a device may register to an account before project assignment) both NOT NULL after backfill (artifacts are uploaded within a project)	backfill account; project_id may stay NULL
Artifact (artifacts :33)	yes	yes	both NOT NULL	backfill account + default project
Release (releases :47)	yes	yes	both NOT NULL	scopes the monotonicity key (\$2.2 #3)
Deployment (deployments :61)	yes	yes	both NOT NULL	scopes the active- uniqueness key (\$ #4)
TelemetryRecord (telemetry_events :73)	yes	inherit via device/ deployment	account_id NOT NULL ; project_id NULLABLE	derive from the referenced device/ deployment at write time
Group (device_groups :96)	yes	yes	account_id NOT NULL ; project_id NULLABLE (a group may be	name becomes unique-per-account (\$2.2 #6)

Entity (struct / table)	add account_id	add project_id	Nullability	Migration / backf implication
			account- wide)	
AuditEntry (audit_logs :115; store.go:123-134)	yes	yes	both NULLABLE	populate from the resolved account/ project context; al finally populate th always-empty User (10_* §6, audit_wire.go:41-

2.2 The 7 single-global-tenant assumptions (10_* §8) → account-scoped

Each global uniqueness/latest key becomes account- (and where relevant project-) scoped:

1. **Project name globally unique** (schema_postgres.sql:241; memory.go:39 prjByName) → **UNIQUE (account_id, name)**. Two tenants may both have a project “production”.
2. **Device hardware_id globally unique** (schema_postgres.sql:30; store.go:338 GetDeviceByHardwareID) → **UNIQUE (account_id, hardware_id)**, and GetDeviceByHardwareID gains an accountID parameter (§3.2). *Tradeoff/flag*: if a single physical device could legitimately belong to two accounts this must instead be a global unique — but the mandate’s tree makes a device account-owned, so per-account is recommended; surfaced for operator confirmation.
3. **Release monotonicity / “latest” keyed (os_type, target_model)** (store.go:349 LatestRelease; releaseMu invariant server.go:56-69) → **(account_id, project_id, os_type, target_model)**. LatestRelease / ReleaseByVersion gain account+project scope (§3.2).
4. **Active-deployment uniqueness keyed (os, target_model, group)** (store.go:359 ActiveDeploymentForTarget; deployMu server.go:42-54) → **(account_id, project_id, os, target_model, group)**.
5. **ProjectAccess.CallerID a bare subject string** (store.go:307-311) → resolves to user_id (§1.2); the same person in several accounts is now representable because membership (§1.3) is keyed per account. project_access rows are implicitly account-scoped via projects.account_id.
6. **Group name globally unique** (schema_postgres.sql:102; memory.go grpByName) → **UNIQUE (account_id, name)**.

7. **Token claims carry no tenant dimension** (token.go:31-36) → the data model requires either an account_id claim in the minted token OR a per-request membership lookup. **The token-shape choice is 21_***s (OQ-2). The *data-model requirement* this doc fixes: the Membership table MUST support a fast ListUserAccounts(userID) (for account selection after sign-in, 00_INDEX §1.5) and GetAccountMembership(userID, accountID) (per-request authZ) — both indexed in §4. **Recommendation to 21_* (not decided here):** resolve membership per-request (supports account switching without re-minting) with an optional account_id claim stamped *after* the user selects an account; tradeoff is a per-request lookup vs a re-mint on switch.

3. store.Repository extension

The Repository is one flat interface (store.go:315-429) with two implementations — MemoryRepository (default, main.go:80) and PostgresRepository (HELIX_DATABASE_URL set, main.go:46-78). The existing Projects block (store.go:317-332) is the exact template. The SAME interface methods are added once and implemented by BOTH backends (10_* §5: “additive-friendly — yes”).

3.1 New interface methods (mirror the Projects block)

```
// Accounts (mirror CreateProject/Get/List/Update/Delete,
store.go:317-321).
CreateAccount(ctx context.Context, a Account) error
GetAccount(ctx context.Context, accountID string) (Account, error)
GetAccountBySlug(ctx context.Context, slug string) (Account,
error) // stable-name resolve (§11.4.111)
ListAccounts(ctx context.Context) ([]Account,
error) // super-admin only (§6)
UpdateAccount(ctx context.Context, a Account) error
DeleteAccount(ctx context.Context, accountID string) error

// Users – the CRUD that does NOT exist today (10_* §3).
CreateUser(ctx context.Context, u User) error
GetUser(ctx context.Context, userID string) (User, error)
GetUserByUsername(ctx context.Context, username string) (User,
error) // token Subject → user
GetUserByEmail(ctx context.Context, email string) (User, error)
ListUsers(ctx context.Context) ([]User,
error) // super-admin only
UpdateUser(ctx context.Context, u User) error
```

```

DeleteUser(ctx context.Context, userID string) error

// Memberships (mirror the ProjectAccess block, store.go:322-332).
SetAccountMembership(ctx context.Context, m AccountMembership)
    error // grant/update role
GetAccountMembership(ctx context.Context, userID, accountID
    string) (AccountMembership, error)
ListAccountMembers(ctx context.Context, accountID string)
    ([]AccountMembership, error)
ListUserAccounts(ctx context.Context, userID string)
    ([]AccountMembership, error) // account selection after
    sign-in
RemoveAccountMembership(ctx context.Context, userID, accountID
    string) error

```

3.2 Scoping the existing read/write methods

Every existing OTA accessor must become account-scoped. Two mechanisms, with a recommended hybrid:

- **Option A — explicit accountID parameter on every scoped method.** e.g. `GetDevice(ctx, accountID, deviceID)`, `GetDeviceByHardwareID(ctx, accountID, hardwareID)`, `LatestRelease(ctx, accountID, projectID, os, targetModel)`, `ActiveDeploymentForTarget(ctx, accountID, projectID, os, targetModel, group)`. *Pro*: the compiler forces every call site to supply scope — a forgotten scope is a build error, not a runtime leak. *Con*: churns every signature + call site.
- **Option B — carry scope in context.Context** (a `ScopeContext{AccountID, ProjectID}` read by each method). *Pro*: minimal signature churn; mirrors the RLS session-variable pattern (set once per request). *Con*: scope becomes implicit — a call in a context missing the scope silently reads unscoped (the exact leak RLS must then catch).
- **Existing *Filter structs already exist** for the list methods (`DeviceFilter` `store.go:280-285`, `ReleaseFilter`, `AuditFilter`) — add `AccountID + ProjectID` fields there for the list paths with zero signature change.

Recommendation (hybrid): Option A (explicit `accountID`, and `projectId` where the resource is project-owned) for the single-entity get/create/update methods + extend the `*Filter` structs for the list methods, with `pgx` RLS as the independent second layer so even a missed scope cannot leak (defense in depth). Explicit params are chosen over `ctx`-implicit specifically because the whole point of this feature is tenant isolation, and a compile-time guarantee beats a

runtime one. *Tradeoff*: larger diff (every OTA handler + call site touched) — acceptable, and 22_api_surface.md already inventories those call sites.

3.3 The missing-transaction problem (atomic “create account + first super-admin”)

The seam has **no transaction primitive** — the code explicitly does a best-effort, non-atomic create-then-grant for projects (10_* \$5; handlers_project.go:137-154). For accounts this is worse: a crash between “account created” and “owner membership granted” leaves an **ownerless orphan tenant** — nobody can administer it and no membership gates it (an isolation + usability defect). Three seam options:

- **Option A — a purpose-built composite method**

CreateAccountWithOwner(ctx, a Account, ownerUserID string, role AccountRole) error, implemented atomically by each backend (pgx: BEGIN ... COMMIT; memory: a single mutex-guarded critical section that inserts both the account and the owner membership, rolling back the in-memory map insert on any error). *Pro*: smallest seam change; matches the existing pattern of purpose-built methods; solves the one flow that actually needs atomicity now. *Con*: a bespoke method per atomic flow (doesn’t generalize).

- **Option B — a generic transaction primitive** WithTx(ctx, func(RepositoryTx) error) error on the interface (pgx: a real transaction; memory: a global-lock stand-in). *Pro*: general — any future multi-write flow (create project + grant, deploy + audit) becomes atomic. *Con*: large seam change; the memory impl’s “transaction” is only a coarse lock (no true rollback of partial map mutations without extra bookkeeping); bigger blast radius.

- **Option C — keep best-effort compensation** (create account → grant → on grant failure, delete account). *Pro*: zero seam change. *Con*: non-atomic across a crash (the orphan-tenant window remains) — rejected for the security-sensitive bootstrap.

Recommendation: Option A now (CreateAccountWithOwner, the one flow that must be atomic for correctness), and **adopt Option B (WithTx) later** when a second multi-write atomic flow appears — deferred to the implementation phase, not required for MVP. *Tradeoff*: Option A doesn’t generalize, but it is the minimal change that closes the orphan-tenant hole; Option B is the right long-term seam but is over-scoped for a single bootstrap flow today.

4. SQL schema sketch (migration-style DDL)

A store-schema-style migration (opaque TEXT ids, mirroring `schema_postgres.sql`) — call it `002_accounts_multitenancy.up.sql`, additive to the shipped store. The canonical UUID variant would land the same tables in the 001-style schema (`001_initial_schema.up.sql`); the mapping is TEXT ↔ UUID, users/api_keys here match `001_initial_schema.up.sql:39-71`. RLS statements are the recommended second isolation layer (§0). This is a **sketch for the design; not applied**.

```
-- 002_accounts_multitenancy (UP) – additive to
    schema_postgres.sql. TEXT ids to
-- match the shipped store; UUID in the canonical 001-style
    schema. NOT executed.
SET search_path = helix_ota, public;

-- 1) New top-level tenant.
CREATE TABLE IF NOT EXISTS helix_ota.accounts (
    account_id TEXT PRIMARY KEY,
    name       TEXT      NOT NULL UNIQUE,
    slug       TEXT      NOT NULL UNIQUE,
    status     TEXT      NOT NULL DEFAULT 'active'
                CHECK (status IN
                    ('active', 'suspended', 'archived')),
    created_at TIMESTAMPTZ NOT NULL DEFAULT now(),
    updated_at TIMESTAMPTZ NOT NULL DEFAULT now()
);

-- 2) Persisted User (closes the 10_* §3 gap; mirrors
    001_initial_schema users).
CREATE TABLE IF NOT EXISTS helix_ota.users (
    user_id      TEXT PRIMARY KEY,
    username     TEXT      NOT NULL UNIQUE,
    email        TEXT      NOT NULL UNIQUE,
    password_hash TEXT      NOT NULL,
    is_super_admin BOOLEAN NOT NULL DEFAULT FALSE, -- §6
                global bypass flag
    is_active    BOOLEAN NOT NULL DEFAULT TRUE,
    created_at   TIMESTAMPTZ NOT NULL DEFAULT now(),
    updated_at   TIMESTAMPTZ NOT NULL DEFAULT now()
);

-- 2b) api_keys (from 001_initial_schema.up.sql:57-71; non-
    interactive CI/CLI creds).
CREATE TABLE IF NOT EXISTS helix_ota.api_keys (
```

```

api_key_id TEXT PRIMARY KEY,
user_id TEXT NOT NULL REFERENCES
    helix_ota.users(user_id) ON DELETE CASCADE,
account_id TEXT NOT NULL REFERENCES
    helix_ota.accounts(account_id) ON DELETE CASCADE,
key_hash TEXT NOT NULL UNIQUE, -- hash
    only; cleartext shown once (§11.4.10)
name TEXT NOT NULL,
permissions JSONB NOT NULL DEFAULT '{}'::jsonb,
expires_at TIMESTAMPTZ,
revoked_at TIMESTAMPTZ,
created_at TIMESTAMPTZ NOT NULL DEFAULT now()
);
CREATE INDEX IF NOT EXISTS idx_api_keys_user ON
    helix_ota.api_keys (user_id);
CREATE INDEX IF NOT EXISTS idx_api_keys_account ON
    helix_ota.api_keys (account_id);

-- 3) Membership (user ↔ account, M:N) – mirrors project_access
    one level up.
CREATE TABLE IF NOT EXISTS helix_ota.account_members (
    user_id TEXT NOT NULL REFERENCES
        helix_ota.users(user_id) ON DELETE CASCADE,
    account_id TEXT NOT NULL REFERENCES
        helix_ota.accounts(account_id) ON DELETE CASCADE,
    role TEXT NOT NULL DEFAULT 'viewer'
        CHECK (role IN ('viewer', 'operator', 'admin')),
    is_owner BOOLEAN NOT NULL DEFAULT FALSE,
    granted_at TIMESTAMPTZ NOT NULL DEFAULT now(),
    granted_by TEXT NOT NULL DEFAULT '',
    PRIMARY KEY (user_id, account_id)
);
CREATE INDEX IF NOT EXISTS idx_account_members_account ON
    helix_ota.account_members (account_id);
CREATE INDEX IF NOT EXISTS idx_account_members_user ON
    helix_ota.account_members (user_id);

-- 3b) RESERVED "max-flexibility" RBAC (only if 21_* chooses it –
    additive, not now):
-- CREATE TABLE helix_ota.roles (role_id TEXT PK, account_id TEXT
    NOT NULL
-- REFERENCES accounts(account_id) ON DELETE CASCADE, name TEXT
    NOT NULL,
-- UNIQUE (account_id, name));
-- CREATE TABLE helix_ota.permissions (permission_id TEXT PK,
    action TEXT NOT NULL,
-- resource TEXT NOT NULL, UNIQUE (action, resource));
-- CREATE TABLE helix_ota.role_permissions (role_id TEXT,
    permission_id TEXT,
-- PRIMARY KEY (role_id, permission_id));

```

```

-- and account_members.role becomes a role_id FK.

-- 4) Scope the existing OTA tables. account_id nullable now →
    backfill → NOT NULL
--    in 003 (see §5 Option A). project_id added to the resources
    per §2.1.
ALTER TABLE helix_ota.projects          ADD COLUMN IF NOT EXISTS
    account_id TEXT;
ALTER TABLE helix_ota.devices           ADD COLUMN IF NOT EXISTS
    account_id TEXT;
ALTER TABLE helix_ota.devices           ADD COLUMN IF NOT EXISTS
    project_id TEXT; -- NULL-until-assigned
ALTER TABLE helix_ota.artifacts         ADD COLUMN IF NOT EXISTS
    account_id TEXT;
ALTER TABLE helix_ota.artifacts         ADD COLUMN IF NOT EXISTS
    project_id TEXT;
ALTER TABLE helix_ota.releases          ADD COLUMN IF NOT EXISTS
    account_id TEXT;
ALTER TABLE helix_ota.releases          ADD COLUMN IF NOT EXISTS
    project_id TEXT;
ALTER TABLE helix_ota.deployments       ADD COLUMN IF NOT EXISTS
    account_id TEXT;
ALTER TABLE helix_ota.deployments       ADD COLUMN IF NOT EXISTS
    project_id TEXT;
ALTER TABLE helix_ota.device_groups     ADD COLUMN IF NOT EXISTS
    account_id TEXT;
ALTER TABLE helix_ota.device_groups     ADD COLUMN IF NOT EXISTS
    project_id TEXT; -- nullable
ALTER TABLE helix_ota.telemetry_events  ADD COLUMN IF NOT EXISTS
    account_id TEXT;
ALTER TABLE helix_ota.telemetry_events  ADD COLUMN IF NOT EXISTS
    project_id TEXT; -- nullable
ALTER TABLE helix_ota.audit_logs        ADD COLUMN IF NOT EXISTS
    account_id TEXT; -- nullable
ALTER TABLE helix_ota.audit_logs        ADD COLUMN IF NOT EXISTS
    project_id TEXT; -- nullable

-- 5) Composite-FK enforcement so resource.account_id ==
    project.account_id (§2).
--    Requires a UNIQUE (project_id, account_id) on projects to be
    the FK target.
--    (Deferred to 003 after backfill sets account_id, since
    project_id is the PK
--    and account_id starts nullable.)
-- ALTER TABLE helix_ota.projects  ADD CONSTRAINT
    projects_pid_aid_uniq UNIQUE (project_id, account_id);
-- ALTER TABLE helix_ota.artifacts ADD CONSTRAINT
    artifacts_proj_fk
--    FOREIGN KEY (project_id, account_id) REFERENCES
    helix_ota.projects (project_id, account_id);
-- (same composite FK for releases, deployments; devices/groups/
    telemetry where project_id NOT NULL)

```

```

-- 6) Uniqueness re-scoping (§2.2). Old global uniques are
    dropped, per-account added
--    in 003 after backfill (a per-account unique cannot be added
    while account_id is
--    still NULL for legacy rows).
-- ALTER TABLE helix_ota.projects      DROP CONSTRAINT ... ; ADD
    UNIQUE (account_id, name);
-- ALTER TABLE helix_ota.devices      DROP CONSTRAINT
    devices_hardware_id_uniq;
-- ADD CONSTRAINT devices_acct_hw_uniq UNIQUE (account_id,
    hardware_id);
-- ALTER TABLE helix_ota.device_groups DROP CONSTRAINT
    device_groups_name_uniq;
-- ADD CONSTRAINT groups_acct_name_uniq UNIQUE (account_id,
    name);

-- 7) Indexes: account_id as the LEADING column of each hot
    scoping index (AWS RLS
--    indexing rule). Added in 003 once account_id is populated +
    NOT NULL.
-- CREATE INDEX idx_devices_account      ON helix_ota.devices
    (account_id);
-- CREATE INDEX idx_artifacts_account    ON helix_ota.artifacts
    (account_id);
-- CREATE INDEX idx_releases_account     ON helix_ota.releases
    (account_id, project_id, os_type, target_model);
-- CREATE INDEX idx_deployments_account  ON helix_ota.deployments
    (account_id, project_id, status);
-- CREATE INDEX idx_telemetry_account    ON
    helix_ota.telemetry_events (account_id);
-- CREATE INDEX idx_audit_account        ON helix_ota.audit_logs
    (account_id);

-- 8) RLS (recommended §0 second layer; enable in 003 after NOT
    NULL). Per-table:
-- ALTER TABLE helix_ota.devices ENABLE ROW LEVEL SECURITY;
-- CREATE POLICY devices_tenant_isolation ON helix_ota.devices
--    USING      (account_id =
    current_setting('app.current_account', true))
--    WITH CHECK (account_id =
    current_setting('app.current_account', true));
-- The app connects as a NON-owner, NON-BYPASSRLS role and runs
-- `SET app.current_account = '<account_id>'` per request (AWS
    pattern). A
-- super-admin request either sets a super-admin GUC the policy
    also honors, or
-- is served by a policy-based bypass (§6) – NOT a BYPASSRLS
    role.

```

Reference to the canonical target schema. users + api_keys above mirror 001_initial_schema.up.sql:39-71 (which the shipped store never carried, 10_* §3). The account_id/project_id

scoping columns are the additive complement to that migration's device/artifact/release/deployment tables (001_initial_schema.up.sql:80-360). When the project adopts the 001-style normalized schema, this becomes 002 there too, with UUID ids and FK-to-users ON uploaded_by/published_by/created_by (001_initial_schema.up.sql:155,219,265) finally resolvable.

5. OQ-4 — migrating existing data under accounts

OQ-4 (00_INDEX §5): migrate existing projects/rollouts/artifacts under a **default account**, or keep the schema **additive-nullable until backfill**? Two options:

- **Option A — default-account backfill (big-bang, then NOT NULL).** 002 adds `account_id` nullable → a one-shot backfill assigns every existing project, device, artifact, release, deployment, group, telemetry, audit row to a single seeded **__default__ system account** (and a `__default__` project where a `project_id` is required) → 003 sets `account_id` NOT NULL, adds the per-account uniques, the composite FKs, the leading-column indexes, and enables RLS. *Pro*: one clean invariant (`account_id` is NEVER NULL for a tenant row after 003); RLS is correct immediately; no “NULL means legacy-global” special case litters the code. *Con*: requires a data-migration step and a brief coordinated 003 cutover.
- **Option B — additive-nullable-until-backfill (gradual).** Add `account_id` nullable and ship code that tolerates NULL as “unassigned/legacy-global”, backfill lazily, flip NOT NULL “eventually.” *Pro*: purely additive, zero-downtime, no cutover. *Con*: while `account_id` is NULL, those rows are **outside** per-account scoping — an isolation hole for the whole window; every read path needs a NULL branch; the NOT-NULL flip still needs the same backfill Option A does, just later and with more accumulated NULL rows.

Recommendation: Option A (default-account backfill). It removes the nullable- isolation-hole window entirely and yields a clean NOT-NULL + RLS invariant on day one. The deciding factor is **data volume**, and the as-is state makes the backfill nearly free: the store is an **in-memory MVP by default** (main.go:80) with Postgres wired only when `HELIX_DATABASE_URL` is set and barely populated — so the big-bang cost Option A normally carries is negligible here.

Operator decision point (flagged, not decided): if, contrary to the as-is state, a target deployment already holds meaningful production OTA data AND a zero-downtime cutover is mandatory, choose Option B and accept the interim NULL- scoping window with an explicit “NULL account_id is visible only to super-admin” guard. The recommendation flips on that single fact — surfaced per §11.4.6, not assumed.

6. Super-admin at the data layer

The mandate: the super-user “sees and controls everything” and administers all accounts/users/permissions (00_INDEX §1.4, §1.6). Model it as a **global users.is_super_admin boolean** (§1.2) — a property of the *identity*, NOT a per-account membership role (a super-admin need not be a member of any account).

How the flag bypasses account scoping (two layers):

- 1. Application layer.** The scope resolver returns “all accounts” when the authenticated user is super-admin — the *exact precedent already in the code*: `requireProjectAccess` lets a global admin bypass the per-project ACL and implicitly access every project (`handlers_project.go:43-60`), and the list-all path returns every project for a global admin (`handlers_project.go:162-199`). The account layer generalizes this: super-admin = skip the `account_id` filter / `requireAccountAccess`.
- 2. Database layer (RLS).** Prefer a **policy-based bypass over a BYPASSRLS role**. A normal request connects as a non-owner, non-BYPASSRLS role and sets `app.current_account`; a super-admin request additionally sets a super-admin GUC (e.g. `SET app.is_super_admin = 'on'`) that each policy’s USING clause also accepts: `USING (account_id = current_setting('app.current_account', true) OR current_setting('app.is_super_admin', true) = 'on')`. AWS explicitly recommends “policy-based admin access over privilege-based bypass” because it “keeps admin access **auditable and revocable**” — a BYPASSRLS role is invisible to the policy and cannot be scoped or logged per query. (PostgreSQL: superusers and BYPASSRLS roles *always* bypass RLS — which is exactly why the app must NOT connect as one for super-admin requests.)

Security implications (data layer; full model in 26_security_threat_model.md):

- **Maximum blast radius.** One compromised super-admin credential = total cross-tenant breach. Therefore `is_super_admin` is settable **only via config/env bootstrap, never a request** — the same trust-boundary rule the codebase already enforces for the token secret (config.go:180-184), the admin bootstrap (main.go:96-104), the TLS-proxy trust flag (config.go:88-110), and the artifact-signature key `resolvePublicKey` (handlers_artifact.go:274-283, which takes the verify key ONLY from server config, never the request). The first super-admin is seeded from a `HELIX_*` env field (or the existing `HELIX_ADMIN_PASSWORD` path extended with `is_super_admin=true`), following `10_* §7` exactly.
 - **Auditability.** Every super-admin action MUST record the **affected tenant's account_id** in the audit row (§2.1; the audit table today carries no `account_id` — `10_* §6`). The policy-based bypass (not a `BYPASSRLS` role) is what keeps those actions attributable per query.
 - **Testable, revocable (anti-bluff).** `26_*` must carry a paired mutation proving
 1. a non-super-admin session with `app.current_account = A` **cannot** read account B's rows even with a forged `WHERE`, and (b) clearing `is_super_admin` (or the GUC) **immediately** re-scopes the session — a super-admin bypass that survives revocation is a §11.4 isolation-layer bluff.
-

Sources verified 2026-07-10

- PostgreSQL Row-Level Security for multi-tenant isolation — runtime tenant context via `SET/ current_setting( 'app.current_tenant' )`, the app-role-must-not-be-owner- or-BYPASSRLS caveat, and the “prefer policy-based admin access over privilege-based bypass ... auditable and revocable” guidance: <https://aws.amazon.com/blogs/database/multi-tenant-data-isolation-with-postgresql-row-level-security/>
- PostgreSQL official docs — canonical RLS primitives (`ALTER TABLE ... ENABLE ROW LEVEL SECURITY, CREATE POLICY ... USING ... WITH CHECK, FORCE ROW LEVEL SECURITY, superuser/BYPASSRLS bypass`): <https://www.postgresql.org/docs/current/ddl-rowsecurity.html>
- WorkOS — multi-tenant RBAC data model: membership join keyed (`user_id, tenant_id, role_id`), global vs tenant-scoped vs hybrid/template roles, `UNIQUE(tenant_id, name)` for tenant-scoped role

namespaces: <https://workos.com/blog/how-to-design-multi-tenant-rbac-saas>

- Microsoft Azure Architecture Center — tenancy models (shared vs schema-per vs database-per) and the isolation/ops-cost tradeoff taxonomy: <https://learn.microsoft.com/en-us/azure/architecture/guide/multitenant/considerations/tenancy-models> and Azure SQL multitenant SaaS patterns: <https://learn.microsoft.com/en-us/azure/azure-sql/database/saas-tenancy-app-design-patterns>
- Supplementary comparison of the three tenancy patterns (shared-schema+RLS as the default recommendation): dasroot 2026, <https://dasroot.net/posts/2026/01/multi-tenancy-database-patterns-schema-database-row-level/>

The composite-FK enforcement of `resource.account_id == project.account_id` (§2) and the `CreateAccountWithOwner` atomic-bootstrap seam (§3.3) are **original work** applied to this schema — no external source prescribes them; they follow from the shipped store’s shape.

Honest boundary (§11.4.6)

- **This is a design proposal, not implemented.** No entity, column, interface method, migration, or RLS policy described here exists in the codebase yet; every “as-is” fact carries a `file:line` (via `00_/10_*`), every “to-be” element is a proposal. Nothing here is a §11.4 completion claim.
- **Every open choice is presented with a recommendation + tradeoffs for operator decision, never silently decided:** the tenancy model (§0), the id type TEXT-vs-UUID divergence (§1), the OTA-resource scoping shape (§2), the Repository scoping mechanism and the atomic-bootstrap seam (§3), the OQ-4 migration path (§5), the super-admin RLS-bypass mechanism (§6). The permission-model *shape* (OQ-2) is deliberately deferred to `21_authz_rbac_superadmin.md`; this doc fixes only the minimal role skeleton + the data-layer requirements `21_*` must satisfy.
- **Scope of the audit behind the as-is facts:** the `file:line` references were established by `10_*` against non-test source only; a test path may exercise a detail not seen. The pgx read paths were read as SQL strings, not executed against a live DB (`10_*` “Honest gaps”). The rollout subsystem’s tenancy (`internal/rollout/`) is marked UNCONFIRMED: in `10_*` and is not designed here — `21_*/30_*` must audit `rollout/store.go` before scoping it.

- **What the recommendation does NOT prove:** shared-schema+RLS is the correct *default*, not a guarantee of isolation — isolation is only as good as the policies
 - the scope-resolution code, which is precisely why 26_security_threat_model.md owns the paired-mutation proof that cross-account reads are impossible.