

SYSTEM_ARCHITECTURE

Helix OTA — System Architecture Document

Document ID: HELOTA-ARCH-001 **Version:** 1.0.0 **Status:** Active **Last Updated:** 2026-03-05 **Constitution Reference:** HelixConstitution v1 §1-§4 **Target Platform:** Android 15 on Orange Pi 5 Max (RK3588)

Table of Contents

- [1. Executive Summary](#)
- [2. High-Level Architecture](#)
- [3. Detailed Component Design](#)
- [4. Data Flow Diagrams](#)
- [5. Android 15 / RK3588 Specific Architecture](#)
- [6. Integration with Existing Submodules](#)
- [7. Communication Protocols](#)
- [8. Security Architecture](#)
- [9. Scalability Design](#)

1. Executive Summary

1.1 System Purpose

Helix OTA is an enterprise-grade Over-The-Air update system designed to deliver secure, reliable, and auditable software updates to fleets of devices. The initial MVP release (1.0.0) targets **Android 15 running on Orange Pi 5 Max (RK3588)**, providing a complete update lifecycle: device registration, update checking, artifact download, cryptographic verification, A/B partition update application, and status reporting.

The system exists because no existing open-source OTA solution provides the combination of: (1) native Android `update_engine` integration, (2) RK3588-specific partition and boot control support, (3) enterprise phased rollout with deterministic cohort assignment, and (4) a Go-based server architecture with first-class observability. Existing solutions like Mender target Linux, not Android's

Virtual A/B mechanism. hawkBit provides fleet management but is Java/Spring-based and client-agnostic. Google's Omaha infrastructure is proprietary and unavailable for self-hosting. Helix OTA fills this gap.

1.2 Design Philosophy

The architecture is governed by four foundational principles:

Principle	Definition	Architectural Manifestation
Universal	The system must eventually support any OS target without architectural changes	OS Adapter interface abstracts platform-specific logic; the 1.0.0 Android adapter is the first implementation of this interface, not a special case
Generic	No device-specific or customer-specific logic in core services	Device model, partition layout, and update mechanism are configuration-driven, not hard-coded; rk3588_opi5max is a device profile, not a code path
Decoupled	Components communicate through well-defined interfaces and events, never through direct coupling	Services communicate via EventBus for async operations and repository interfaces for data access; no service directly imports another service's implementation
Extensible	New capabilities are added by implementing interfaces, not by modifying existing code	Storage backends (S3, MinIO, local), authentication methods (mTLS, JWT), and update mechanisms (full, delta) are all pluggable via interface registration

These principles ensure that the system evolves from an Android-only MVP (1.0.0) to a universal multi-OS platform (2.0.0) through additive extension, never through architectural rewrite.


1.3 Scope Boundaries


In Scope (1.0.0-MVP): - Full A/B update pipeline for Android 15 on RK3588 - REST API server in Go with phased rollout engine - Android client integrating with update_engine via Binder IPC - Web dashboard for artifact upload, rollout management, and fleet monitoring - RSA-4096 artifact signing and SHA-256 verification - mTLS device authentication and JWT dashboard authentication - PostgreSQL 16 for all persistent storage - S3-compatible artifact storage (MinIO or AWS S3)

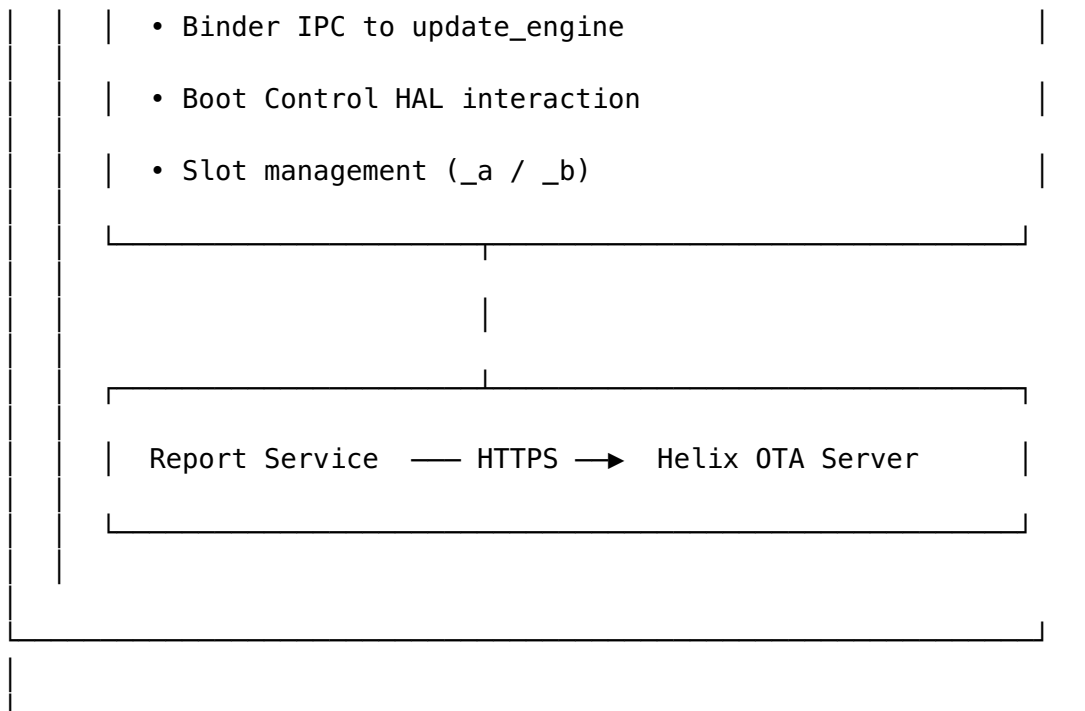
Out of Scope (1.0.0-MVP): - Delta/differential updates (1.0.2) - Multi-version rollback beyond A/B slot swap (1.0.1) - Linux and Windows support (1.1.0, 1.2.0) - Plugin architecture (2.0.0) - Google Omaha protocol compatibility

2. High-Level Architecture

2.1 System Topology







2.2 Five Major Subsystems

#	Subsystem	Technology	Primary Responsibility
1	Helix OTA Server	Go 1.22+, Chi router, gRPC	Central update management: artifact validation, device registry, rollout engine, telemetry ingestion, authentication
2	Helix OTA Client SDK	Go library (cross-compiled for Android via gomobile)	Reusable update client logic: check, download, verify, install orchestration, status reporting
3	Helix OTA Android Client	Kotlin/Java (AOSP integration layer)	Android-specific integration: Binder IPC to update_engine, boot control HAL, WorkManager scheduling
4	Helix OTA Dashboard	React 18+, TypeScript, Vite	Administrative web UI: artifact upload, rollout management, fleet monitoring, telemetry visualization
5	Helix OTA Infrastructure	Docker Compose (dev), Kubernetes (prod), PostgreSQL 16, Redis 7, MinIO	Deployment and data persistence: containers, databases, message queues, object storage

2.3 Subsystem Interaction Matrix

	Server Infrastructure	Client SDK	Android Client
Dashboard Server	–	provides	integrates

serves	depends on		
Client SDK	embedded	–	wraps
–	configures for		
Android Client	calls	uses	
–	–	reads from	
Dashboard	reads		
–	–	–	renders
Infrastructure	hosts		
–	–	–	–

3. Detailed Component Design

3.1 Server Components

The Helix OTA Server is a monolithic Go application structured as a set of loosely-coupled service modules, each owning a bounded context. Services communicate synchronously through repository interfaces and asynchronously through the EventBus submodule.

3.1.1 Update Service

Responsibility: Manages update artifacts, version compatibility, and update availability decisions.

```
// UpdateService manages the update lifecycle from the server's
// perspective.
// It determines what update is available for a given device based
// on:
//   - Device's current version
//   - Active rollouts targeting the device's group
//   - Rollout cohort assignment (deterministic hash-based)
type UpdateService struct {
    deviceRepo DeviceRepository
    artifactRepo ArtifactRepository
    rolloutRepo RolloutRepository
    cache       cache.Provider // vasic-digital/cache
    events      eventbus.Publisher // vasic-digital/EventBus
}

// CheckForUpdate determines whether an update is available for a
// device.
// The rollout cohort is computed as: cohort = fnv32(device_id) %
// 100
// A device receives the update if cohort <
// rollout.current_percentage.
func (s *UpdateService) CheckForUpdate(
    ctx context.Context,
    deviceID string,
) (*UpdateInfo, error) {
    // 1. Load device from cache or database
    device, err := s.deviceRepo.GetByID(ctx, deviceID)
```

```

    if err != nil {
        return nil, fmt.Errorf("device lookup: %w", err)
    }

    // 2. Find active rollout targeting this device's group
    rollout, err := s.rolloutRepo.FindActiveForGroup(ctx,
        device.Group)
    if err != nil {
        return nil, fmt.Errorf("rollout lookup: %w", err)
    }
    if rollout == nil {
        return nil, nil // no update available
    }

    // 3. Deterministic cohort assignment
    cohort := fnv32Hash(deviceID) % 100
    if cohort >= rollout.CurrentPercentage {
        return nil, nil // not in rollout cohort yet
    }

    // 4. Check version compatibility
    artifact, err := s.artifactRepo.GetByID(ctx,
        rollout.ArtifactID)
    if err != nil {
        return nil, fmt.Errorf("artifact lookup: %w", err)
    }
    if !isCompatible(device.CurrentVersion,
        artifact.TargetVersion, artifact.MinSourceVersion) {
        return nil, nil // version not compatible
    }

    return &UpdateInfo{
        ArtifactID:  artifact.ID,
        Version:     artifact.TargetVersion,
        DownloadURL: fmt.Sprintf("/api/v1/artifacts/%s/download", artifact.ID),
        SHA256:     artifact.SHA256,
        SizeBytes:   artifact.SizeBytes,
        SignatureURL: fmt.Sprintf("/api/v1/artifacts/%s/signature", artifact.ID),
        Metadata: UpdateMetadata{
            Mandatory: rollout.Mandatory,
            Deadline:  rollout.Deadline,
        },
    }, nil
}

// UpdateInfo represents an available update returned to the
// client.
type UpdateInfo struct {
    ArtifactID string `json:"artifact_id"`
    Version    string `json:"version"`
    DownloadURL string `json:"download_url"`
    SHA256     string `json:"sha256"`
}

```

```

SizeBytes      int64      `json:"size_bytes"`
SignatureURL   string     `json:"signature_url"`
Metadata      UpdateMetadata `json:"metadata"`
}

```

Key Design Decisions: - Cohort assignment uses FNV-32 hash for deterministic, stable device-to-cohort mapping. A device assigned to the 5% cohort remains in the 5% cohort even if the rollout later increases to 30% — it never “goes back” to no-update. - Version compatibility uses semver comparison with an explicit MinSourceVersion field on each artifact, allowing non-linear version jumps (e.g., skipping from 15.0.0 to 15.0.2). - Results are cached in Redis with a TTL of 60 seconds to avoid thundering-herd queries on rollout percentage changes.

3.1.2 Device Service

Responsibility: Device registration, inventory management, and device grouping.

```

// DeviceService handles device lifecycle management.
type DeviceService struct {
    repo DeviceRepository
    cache cache.Provider
    events eventbus.Publisher
}

// RegisterDevice creates a new device record and issues
// authentication credentials.
func (s *DeviceService) RegisterDevice(
    ctx context.Context,
    req DeviceRegistrationRequest,
) (*Device, string, error) {
    // 1. Check for duplicate by hardware fingerprint
    existing, err := s.repo.GetByFingerprint(ctx,
        req.HardwareFingerprint)
    if err == nil && existing != nil {
        // Re-registration: rotate credentials, update metadata
        token, err := s.rotateCredentials(ctx, existing.ID)
        return existing, token, err
    }

    // 2. Create device record
    device := &Device{
        ID:                generateID("dev_"),
        Serial:            req.Serial,
        Model:             req.Model,
        CurrentVersion:    req.CurrentVersion,
        SlotSuffix:        req.SlotSuffix,
        HardwareFingerprint: req.HardwareFingerprint,
        Group:             determineGroup(req.Model), // e.g.,
        "rk3588_opi5max"
        Status:           DeviceStatusIdle,
        LastCheckIn:       time.Now(),
    }
}

```



```

    if err := s.repo.Create(ctx, device); err != nil {
        return nil, "", fmt.Errorf("create device: %w", err)
    }

    // 3. Issue JWT device token
    token, err := s.issueDeviceToken(ctx, device.ID)
    if err != nil {
        return nil, "", fmt.Errorf("issue token: %w", err)
    }

    // 4. Publish device registered event
    s.events.Publish(ctx, eventbus.Event{
        Type:    "device.registered",
        Payload: device,
    })

    return device, token, nil
}

// Device represents a registered device in the system.
type Device struct {
    ID          string    `json:"id" db:"id"`
    Serial      string    `json:"serial" db:"serial"`
    Model       string    `json:"model" db:"model"`
    CurrentVersion string  `json:"current_version" db:"current_version"`
    TargetVersion *string  `json:"target_version" db:"target_version"`
    SlotSuffix  string    `json:"slot_suffix" db:"slot_suffix"`
    HardwareFingerprint string `json:"hardware_fingerprint" db:"hardware_fingerprint"`
    Group       string    `json:"device_group" db:"device_group"`
    Status      string    `json:"status" db:"status"`
    LastCheckIn time.Time `json:"last_check_in" db:"last_check_in"`
    CreatedAt   time.Time `json:"created_at" db:"created_at"`
    UpdatedAt   time.Time `json:"updated_at" db:"updated_at"`
}

```

Device Group Strategy: Devices are automatically assigned to groups based on their `model` field. For MVP, the primary group is `rk3588_opi5max`. Group assignment is extensible — future versions will support custom group rules (e.g., by firmware version, region, or customer). Groups determine rollout targeting.

3.1.3 Rollout Service

Responsibility: Phased deployment engine with deterministic cohort assignment, pause/resume, and auto-advance.

```
// RolloutService manages the lifecycle of update rollouts.
type RolloutService struct {
    repo          RolloutRepository
    artifactRepo  ArtifactRepository
    deviceRepo    DeviceRepository
    events        eventbus.Publisher
    workers       *concurrency.Pool // vasic-digital/concurrency
}

// Rollout represents a phased deployment of an artifact to a
// device group.
type Rollout struct {
    ID                string    `json:"id" db:"id"`
    ArtifactID        string    `json:"artifact_id" db:"artifact_id"`
    DeviceGroup        string    `json:"device_group" db:"device_group"`
    CurrentPercentage int       `json:"current_percentage" db:"current_percentage"`
    TargetPercentage  int       `json:"target_percentage" db:"target_percentage"`
    IncrementStep      int       `json:"increment_step" db:"increment_step"`
    AutoAdvance        bool      `json:"auto_advance" db:"auto_advance"`
    Status             string    `json:"status" db:"status" // ACTIVE, PAUSED, COMPLETED, ROLLED_BACK`
    Mandatory          bool      `json:"mandatory" db:"mandatory"`
    Deadline           *time.Time `json:"deadline" db:"deadline"`
    CreatedAt          time.Time `json:"created_at" db:"created_at"`
    UpdatedAt          time.Time `json:"updated_at" db:"updated_at"`
}

// CreateRollout initializes a new phased rollout.
func (s *RolloutService) CreateRollout(
    ctx context.Context,
    req CreateRolloutRequest,
) (*Rollout, error) {
    // Validate artifact exists and is valid
    artifact, err := s.artifactRepo.GetByID(ctx, req.ArtifactID)
    if err != nil {
        return nil, fmt.Errorf("artifact lookup: %w", err)
    }
    if artifact.ValidationStatus != "VALID" {
        return nil, ErrArtifactNotValid
    }
}
```

```

    }

    rollout := &Rollout{
        ID:                generateID("rol_"),
        ArtifactID:        req.ArtifactID,
        DeviceGroup:       req.DeviceGroup,
        CurrentPercentage: req.InitialPercentage,
        TargetPercentage:  100,
        IncrementStep:     req.IncrementStep,
        AutoAdvance:       req.AutoAdvance,
        Status:            "ACTIVE",
        Mandatory:         req.Mandatory,
        Deadline:          req.Deadline,
    }

    if err := s.repo.Create(ctx, rollout); err != nil {
        return nil, fmt.Errorf("create rollout: %w", err)
    }

    s.events.Publish(ctx, eventbus.Event{
        Type:    "rollout.created",
        Payload: rollout,
    })

    return rollout, nil
}

// AdvanceRollout increases the rollout percentage by the
// configured increment step.
// Auto-advance is triggered by a scheduled job that checks
// rollout health metrics.
func (s *RolloutService) AdvanceRollout(
    ctx context.Context,
    rolloutID string,
    targetPercentage int,
) (*Rollout, error) {
    rollout, err := s.repo.GetByID(ctx, rolloutID)
    if err != nil {
        return nil, fmt.Errorf("rollout lookup: %w", err)
    }
    if rollout.Status != "ACTIVE" {
        return nil, ErrRolloutNotActive
    }
    if targetPercentage <= rollout.CurrentPercentage {
        return nil, ErrInvalidPercentage
    }
    if targetPercentage > 100 {
        return nil, ErrPercentageOverflow
    }

    // Health check before advance: ensure failure rate < 5% in
    // current cohort
    healthMetrics, err := s.getRolloutHealth(ctx, rolloutID)

```

```

    if err != nil {
        return nil, fmt.Errorf("health check: %w", err)
    }
    if healthMetrics.FailureRate > 0.05 {
        return nil, ErrHighFailureRate(healthMetrics.FailureRate)
    }

    rollout.CurrentPercentage = targetPercentage
    if targetPercentage >= 100 {
        rollout.Status = "COMPLETED"
    }

    if err := s.repo.Update(ctx, rollout); err != nil {
        return nil, fmt.Errorf("update rollout: %w", err)
    }

    // Invalidate device check-in cache so devices re-query
    s.events.Publish(ctx, eventbus.Event{
        Type:    "rollout.advanced",
        Payload: rollout,
    })

    return rollout, nil
}

```

Rollout Percentage Progression:

Standard: 5% → 10% → 30% → 50% → 100%
 Conservative: 1% → 5% → 10% → 25% → 50% → 100%
 Aggressive: 10% → 25% → 50% → 100%
 Custom: User-defined increments via dashboard

Auto-Advance Logic: When AutoAdvance is enabled, a background goroutine evaluates rollout health every 15 minutes. Advance occurs only if: (1) $\geq 80\%$ of devices in the current cohort have reported a terminal status (SUCCESS or FAILED), (2) failure rate $< 5\%$, and (3) no pause has been manually triggered.

3.1.4 Artifact Service

Responsibility: Upload, multi-stage validation, and storage of OTA artifacts.

```

// ArtifactService manages OTA artifact lifecycle.
type ArtifactService struct {
    repo      ArtifactRepository
    storage    storage.Provider // vasic-digital/storage
    workers    *concurrency.Pool // vasic-digital/concurrency
    verifier   ArtifactVerifier
    events     eventbus.Publisher
    signingKey *rsa.PrivateKey // loaded from HSM/vault in
                  production
}

// Artifact represents an uploaded OTA update package.

```

```

type Artifact struct {
    ID            string    `json:"id"
                        db:"id"`
    Filename      string    `json:"filename"
                        db:"filename"`
    TargetVersion string    `json:"target_version"
                        db:"target_version"`
    TargetModel   string    `json:"target_model"
                        db:"target_model"`
    MinSourceVersion string  `json:"min_source_version"
                        db:"min_source_version"`
    SHA256        string    `json:"sha256"
                        db:"sha256"`
    SizeBytes     int64    `json:"size_bytes"
                        db:"size_bytes"`
    StorageKey    string    `json:"- "
                        db:"storage_key"`
    Signature     []byte   `json:"- "
                        db:"signature"`
    ValidationStatus string  `json:"validation_status"
                        db:"validation_status"`
    CreatedAt     time.Time `json:"created_at"
                        db:"created_at"`
}

// UploadArtifact processes an uploaded OTA zip through the full
// validation chain.
// The validation chain is: hash verification → RSA signature →
// zip structure →
// compatibility check. All four stages must pass for the artifact
// to be marked VALID.
func (s *ArtifactService) UploadArtifact(
    ctx context.Context,
    reader io.Reader,
    metadata ArtifactMetadata,
) (*Artifact, error) {
    // 1. Stream to temporary storage while computing SHA-256
    tempPath, sha256Hash, size, err :=
        s.streamToTempWithHash(ctx, reader)
    if err != nil {
        return nil, fmt.Errorf("stream to temp: %w", err)
    }
    defer os.Remove(tempPath)

    artifact := &Artifact{
        ID:            generateID("art_"),
        Filename:      metadata.Filename,
        TargetVersion: metadata.TargetVersion,
        TargetModel:   metadata.TargetModel,
        SHA256:        sha256Hash,
        SizeBytes:     size,
        ValidationStatus: "PENDING",
    }

    // 2. Run validation chain in worker pool

```

```

result, err := s.workers.Submit(ctx, func()
    (*ValidationResult, error) {
        return s.verifier.Validate(ctx, tempPath, artifact)
    }).Await(ctx)

if err != nil || !result.Valid {
    artifact.ValidationStatus = "INVALID"
    s.repo.Create(ctx, artifact) // persist for audit
    return nil, ErrArtifactValidation(result.Errors)
}

// 3. Sign the artifact
signature, err := s.signArtifact(artifact)
if err != nil {
    return nil, fmt.Errorf("sign artifact: %w", err)
}
artifact.Signature = signature

// 4. Upload to persistent storage (S3/MinIO)
storageKey := fmt.Sprintf("artifacts/%s/%s", artifact.ID,
    metadata.Filename)
if err := s.storage.Upload(ctx, storageKey, tempPath); err !=
    nil {
    return nil, fmt.Errorf("upload to storage: %w", err)
}
artifact.StorageKey = storageKey
artifact.ValidationStatus = "VALID"

// 5. Persist artifact record
if err := s.repo.Create(ctx, artifact); err != nil {
    return nil, fmt.Errorf("persist artifact: %w", err)
}

s.events.Publish(ctx, eventbus.Event{
    Type:    "artifact.uploaded",
    Payload: artifact,
})

return artifact, nil
}

// ArtifactVerifier performs the multi-stage validation chain.
type ArtifactVerifier interface {
    Validate(ctx context.Context, path string, artifact
        *Artifact) (*ValidationResult, error)
}

// ValidationResult captures the outcome of the validation chain.
type ValidationResult struct {
    Valid bool    `json:"valid"`
    Errors []string `json:"errors,omitempty"`
    Stages []ValidationStage `json:"stages"`
}

```

```

type ValidationStage struct {
    Name      string `json:"name"`      // "hash", "signature",
                "structure", "compatibility"
    Passed    bool   `json:"passed"`
    Message   string `json:"message"`
}

```

Validation Chain (four stages, all must pass):

Stage	Check	Failure Mode	Time Budget
1. Hash	SHA-256 of uploaded file matches computed hash	Tampering or corruption during upload	O(n) in file size
2. Signature	RSA-4096 signature on payload.bin verifies against embedded public key	Unsigned or forged artifact	< 5s
3. Structure	ZIP contains payload.bin + payload_properties.txt + care_map.pb; payload.bin header is valid	Non-conforming OTA package format	< 30s for 2GB
4. Compatibility	target_model matches a known device profile; target_version > min_source_version semantically	Artifact targeting wrong device or downgrade attempt	< 1s

3.1.5 Telemetry Service

Responsibility: Device health monitoring, update metrics aggregation, and anomaly detection.

// TelemetryService ingests and analyzes device telemetry data.

```

type TelemetryService struct {
    repo      TelemetryRepository
    events    eventbus.Publisher
    cache     cache.Provider
    workers   *concurrency.Pool
}

```

// TelemetryEvent represents a single telemetry data point from a device.

```

type TelemetryEvent struct {
    ID          string          db:"id"
    DeviceID    string          db:"device_id"
    ArtifactID  string          db:"artifact_id"
}

```

```

EventType      string
    `json:"event_type"          db:"event_type" ` // CHECK,
                                DOWNLOAD_PROGRESS, VERIFY, APPLY, REBOOT, SUCCESS, FAILURE
Status          string
    `json:"status"              db:"status" `
ProgressPercent int
    `json:"progress_percent"     db:"progress_percent" `
ErrorCode       *string
    `json:"error_code"           db:"error_code" `
ErrorMessage    *string
    `json:"error_message"        db:"error_message" `
Metadata        map[string]interface{}
    `json:"metadata"             db:"metadata" `
Timestamp       time.Time
    `json:"timestamp"            db:"timestamp" `
IngestedAt      time.Time
    `json:"ingested_at"          db:"ingested_at" `
}

// IngestTelemetry processes incoming telemetry events from
// devices.
// Events are published to EventBus for async processing by
// consumers:
// - Metrics aggregator: computes success/failure rates per
//   rollout
// - Anomaly detector: flags unexpected patterns (e.g., spike in
//   failures)
// - Notification dispatcher: alerts dashboard of critical
//   events
func (s *TelemetryService) IngestTelemetry(
    ctx context.Context,
    events []TelemetryEvent,
) error {
    // Batch insert into database
    if err := s.repo.BatchCreate(ctx, events); err != nil {
        return fmt.Errorf("batch create telemetry: %w", err)
    }

    // Publish each event for async processing
    for _, event := range events {
        s.events.Publish(ctx, eventbus.Event{
            Type:    fmt.Sprintf("telemetry.%s", event.EventType),
            Payload: event,
        })
    }

    return nil
}

// GetRolloutHealth computes aggregate health metrics for a
// rollout.
func (s *TelemetryService) GetRolloutHealth(
    ctx context.Context,
    rolloutID string,
) (*RolloutHealth, error) {

```



```

// Try cache first (refreshed every 60s by background job)
cached, err := s.cache.Get(ctx, fmt.Sprintf("rollout:health:
    %s", rolloutID))
if err == nil {
    return cached.(*RolloutHealth), nil
}

// Compute from telemetry events
metrics, err := s.repo.AggregateRolloutMetrics(ctx, rolloutID)
if err != nil {
    return nil, fmt.Errorf("aggregate metrics: %w", err)
}

health := &RolloutHealth{
    RolloutID:      rolloutID,
    TotalDevices:   metrics.TotalDevices,
    SuccessCount:   metrics.SuccessCount,
    FailureCount:   metrics.FailureCount,
    InProgressCount: metrics.InProgressCount,
    FailureRate:    float64(metrics.FailureCount) /
        float64(metrics.TotalDevices),
    AverageDuration: metrics.AverageDuration,
    TopErrorCodes:  metrics.TopErrorCodes,
}

// Cache for 60 seconds
s.cache.Set(ctx, fmt.Sprintf("rollout:health:%s", rolloutID),
    health, 60*time.Second)

return health, nil
}

```

Anomaly Detection Rules: - Failure rate > 5% in any cohort triggers auto-pause of the rollout - > 3 devices reporting the same error code within 5 minutes triggers an alert - Device not reporting for > 2× check interval triggers “device stale” notification

3.1.6 Auth Service

Responsibility: JWT token management, mTLS certificate verification, and RBAC enforcement.

```

// AuthService handles authentication and authorization.
type AuthService struct {
    jwtProvider    auth.JWTProvider    // vasic-digital/auth
    certValidator security.CertValidator // vasic-digital/security
    repo           AuthRepository
}

// Role represents a user role in the RBAC system.
type Role string

const (

```

```

RoleAdmin    Role = "admin"    // Full access: manage
                artifacts, rollouts, devices, users
RoleOperator Role = "operator" // Operational access: manage
                rollouts, view devices/artifacts
RoleViewer   Role = "viewer"   // Read-only access: view
                dashboards, device status, telemetry
)

// RBAC permission matrix
var rolePermissions = map[Role][]Permission{
    RoleAdmin: {
        PermArtifactUpload, PermArtifactDelete,
        PermRolloutCreate, PermRolloutUpdate, PermRolloutPause,
        PermRolloutDelete,
        PermDeviceView, PermDeviceDelete,
        PermUserCreate, PermUserUpdate, PermUserDelete,
        PermTelemetryView,
    },
    RoleOperator: {
        PermArtifactUpload, PermArtifactView,
        PermRolloutCreate, PermRolloutUpdate, PermRolloutPause,
        PermDeviceView,
        PermTelemetryView,
    },
    RoleViewer: {
        PermArtifactView,
        PermRolloutView,
        PermDeviceView,
        PermTelemetryView,
    },
}

// AuthenticateDevice verifies device identity via mTLS or JWT.
// mTLS is preferred: the device's client certificate CN must
// match the device ID.
// JWT fallback: token must contain a valid device_id claim.
func (s *AuthService) AuthenticateDevice(
    ctx context.Context,
    cert *x509.Certificate,
    token string,
) (*DeviceIdentity, error) {
    // Try mTLS first
    if cert != nil {
        deviceID := cert.Subject.CommonName
        if !isValidDeviceID(deviceID) {
            return nil, ErrInvalidDeviceCertificate
        }
        return &DeviceIdentity{
            DeviceID: deviceID,
            AuthMethod: "mTLS",
        }, nil
    }
}

```

```

// Fallback to JWT
claims, err := s.jwtProvider.ValidateToken(ctx, token)
if err != nil {
    return nil, ErrInvalidToken
}

deviceID, ok := claims["device_id"].(string)
if !ok {
    return nil, ErrInvalidTokenClaims
}

return &DeviceIdentity{
    DeviceID: deviceID,
    AuthMethod: "JWT",
}, nil
}

```

3.1.7 Notification Service

Responsibility: WebSocket and Server-Sent Events (SSE) push to the dashboard for real-time updates.

```

// NotificationService manages real-time push notifications to
// dashboard clients.
type NotificationService struct {
    hub      *WebSocketHub
    events   eventbus.Subscriber
    sessions map[string]*DashboardSession
    mu       sync.RWMutex
}

// NotificationType categorizes dashboard push events.
type NotificationType string

const (
    NotifDeviceRegistered NotificationType = "device.registered"
    NotifDeviceStatus     NotificationType =
        "device.status_changed"
    NotifRolloutAdvanced NotificationType = "rollout.advanced"
    NotifRolloutPaused   NotificationType = "rollout.paused"
    NotifArtifactUploaded NotificationType = "artifact.uploaded"
    NotifAnomalyDetected NotificationType = "anomaly.detected"
)

// DashboardSession represents an active dashboard WebSocket/SSE
// connection.
type DashboardSession struct {
    ID          string
    UserID      string
    Role        Role
    Conn        *websocket.Conn
    Filters     []string // e.g., ["device_group:rk3588_opi5max"]
    LastPing    time.Time
}

```

```

}

// Start begins consuming EventBus events and pushing to dashboard
sessions.
func (s *NotificationService) Start(ctx context.Context) error {
    ch := s.events.Subscribe(ctx,
        "device.registered",
        "device.status_changed",
        "rollout.advanced",
        "rollout.paused",
        "artifact.uploaded",
        "anomaly.detected",
    )

    for {
        select {
        case event := <-ch:
            s.broadcastToSessions(ctx, event)
        case <-ctx.Done():
            return ctx.Err()
        }
    }
}

```

3.2 Client Components

The Helix OTA Client SDK is a Go library providing the device-side update logic. It is consumed by the Android Client (via gomobile bindings) and can also be embedded in any Go-based client for future OS targets.

3.2.1 Update Checker

Responsibility: Periodic polling with configurable interval, exponential backoff on failure, and respect for device idle/battery constraints.

```

// UpdateChecker polls the OTA server for available updates.
type UpdateChecker struct {
    client      *http.Client
    serverURL   string
    deviceID    string
    interval    time.Duration
    maxInterval time.Duration // cap for exponential backoff
    currentBackoff time.Duration
    checkChan   chan *UpdateInfo
    errChan     chan error
}

// NewUpdateChecker creates a checker with default 4-hour
// interval.
// On failure, backoff follows: 1h → 2h → 4h → 8h → 24h (max).
func NewUpdateChecker(cfg Config) *UpdateChecker {
    return &UpdateChecker{

```

```

        client:      cfg.HTTPClient,
        serverURL:   cfg.ServerURL,
        deviceID:    cfg.DeviceID,
        interval:    cfg.CheckInterval,    // default: 4 *
        time.Hour
        maxInterval: 24 * time.Hour,
        checkChan:   make(chan *UpdateInfo, 1),
        errChan:     make(chan error, 1),
    }
}

// Check performs a single update check against the server.
func (c *UpdateChecker) Check(ctx context.Context) (*UpdateInfo,
    error) {
    url := fmt.Sprintf("%s/api/v1/devices/%s/update-check",
        c.serverURL, c.deviceID)
    req, err := http.NewRequestWithContext(ctx, http.MethodGet,
        url, nil)
    if err != nil {
        return nil, fmt.Errorf("create request: %w", err)
    }

    resp, err := c.client.Do(req)
    if err != nil {
        c.currentBackoff = c.nextBackoff()
        return nil, fmt.Errorf("check request: %w", err)
    }
    defer resp.Body.Close()

    c.currentBackoff = 0 // reset on success

    if resp.StatusCode == http.StatusNoContent {
        return nil, nil // no update available
    }

    var info UpdateInfo
    if err := json.NewDecoder(resp.Body).Decode(&info); err !=
        nil {
        return nil, fmt.Errorf("decode response: %w", err)
    }

    return &info, nil
}

func (c *UpdateChecker) nextBackoff() time.Duration {
    next := c.currentBackoff * 2
    if next == 0 {
        next = time.Hour
    }
    if next > c.maxInterval {
        next = c.maxInterval
    }
}

```

```

    return next
}

```

3.2.2 Download Manager

Responsibility: Resumable chunked downloads with HTTP Range header support, bandwidth throttling, and encrypted at-rest storage of partial downloads.

```

// DownloadManager handles resumable, throttled artifact
// downloads.
type DownloadManager struct {
    client      *http.Client
    storageDir  string
    throttleRate int64 // bytes per second, 0 = unlimited
}

// DownloadResult contains the path and hash of the completed
// download.
type DownloadResult struct {
    FilePath string
    SHA256   string
    Size     int64
}

// Download downloads an OTA artifact with resume support.
// If a partial download exists, it resumes from the last byte.
func (dm *DownloadManager) Download(
    ctx context.Context,
    url string,
    expectedSHA256 string,
    expectedSize int64,
    progressChan chan<- int, // 0-100 progress
) (*DownloadResult, error) {
    filePath := filepath.Join(dm.storageDir, "download.tmp")
    hash := sha256.New()

    // Check for existing partial download
    var offset int64
    if info, err := os.Stat(filePath); err == nil {
        offset = info.Size()
        // Reopen and compute hash of existing portion
        f, _ := os.Open(filePath)
        io.Copy(hash, f)
        f.Close()
    }

    // Create or append to file
    f, err := os.OpenFile(filePath, os.O_CREATE|os.O_WRONLY|
        os.O_APPEND, 0600)
    if err != nil {
        return nil, fmt.Errorf("open file: %w", err)
    }
    defer f.Close()

```

```

// Create request with Range header for resume
req, err := http.NewRequestWithContext(ctx, http.MethodGet,
    url, nil)
if err != nil {
    return nil, fmt.Errorf("create request: %w", err)
}
if offset > 0 {
    req.Header.Set("Range", fmt.Sprintf("bytes=%d-", offset))
}

resp, err := dm.client.Do(req)
if err != nil {
    return nil, fmt.Errorf("download request: %w", err)
}
defer resp.Body.Close()

// Throttled writer: respects bandwidth limit
writer := io.MultiWriter(f, hash)
if dm.throttleRate > 0 {
    writer = newThrottledWriter(writer, dm.throttleRate)
}

// Progress tracking writer
progressWriter := &progressWriter{
    writer:      writer,
    total:       expectedSize,
    downloaded:  offset,
    progressChan: progressChan,
}

if _, err := io.Copy(progressWriter, resp.Body); err != nil {
    return nil, fmt.Errorf("download copy: %w", err)
}

// Verify SHA-256
computedHash := hex.EncodeToString(hash.Sum(nil))
if computedHash != expectedSHA256 {
    os.Remove(filePath)
    return nil, ErrHashMismatch
}

finalPath := filepath.Join(dm.storageDir, "download.bin")
os.Rename(filePath, finalPath)

return &DownloadResult{
    FilePath: finalPath,
    SHA256:   computedHash,
    Size:     expectedSize,
}, nil
}

```

3.2.3 Verification Engine

Responsibility: SHA-256 hash verification, RSA-4096 signature verification, and payload structure validation.

```
// VerificationEngine verifies OTA artifacts before installation.
type VerificationEngine struct {
    trustedPublicKeys []*rsa.PublicKey // embedded in client
    binary
}

// Verify performs the complete verification chain.
func (ve *VerificationEngine) Verify(
    ctx context.Context,
    artifactPath string,
    expectedSHA256 string,
    signature []byte,
) error {
    // Stage 1: SHA-256 hash verification
    if err := ve.verifyHash(artifactPath, expectedSHA256); err !=
        nil {
        return fmt.Errorf("hash verification failed: %w", err)
    }

    // Stage 2: RSA signature verification
    if err := ve.verifySignature(artifactPath, signature); err !=
        nil {
        return fmt.Errorf("signature verification failed: %w",
            err)
    }

    // Stage 3: Payload structure validation
    if err := ve.verifyPayloadStructure(artifactPath); err != nil
    {
        return fmt.Errorf("payload structure invalid: %w", err)
    }

    return nil
}

// verifySignature verifies the RSA-4096 signature against all
// trusted public keys.
// The signature is computed over the SHA-256 hash of payload.bin.
func (ve *VerificationEngine) verifySignature(
    artifactPath string,
    signature []byte,
) error {
    hash, err := computeFileSHA256(artifactPath)
    if err != nil {
        return fmt.Errorf("compute hash: %w", err)
    }

    hashed := sha256.Sum256([]byte(hash))
    for _, pubKey := range ve.trustedPublicKeys {
```



```

        err := rsa.VerifyPKCS1v15(pubKey, crypto.SHA256,
            hashed[:], signature)
        if err == nil {
            return nil // signature verified
        }
    }

    return ErrSignatureVerificationFailed
}

// verifyPayloadStructure checks that the OTA zip contains the
// required files.
func (ve *VerificationEngine) verifyPayloadStructure(artifactPath
    string) error {
    r, err := zip.OpenReader(artifactPath)
    if err != nil {
        return fmt.Errorf("open zip: %w", err)
    }
    defer r.Close()

    required := map[string]bool{
        "payload.bin":      false,
        "payload_properties.txt": false,
        "care_map.pb":      false,
    }

    for _, f := range r.File {
        if _, ok := required[f.Name]; ok {
            required[f.Name] = true
        }
    }

    for name, found := range required {
        if !found {
            return fmt.Errorf("missing required file: %s", name)
        }
    }

    return nil
}

```

3.2.4 Install Orchestrator

Responsibility: Interfaces with Android's update_engine via Binder IPC to apply A/B updates.

```

// InstallOrchestrator drives update_engine to apply A/B partition
// updates.
// On Android, this communicates via Binder IPC to the
// update_engine daemon.
// The Go SDK defines the interface; the Android Client provides
// the Binder implementation.
type InstallOrchestrator interface {

```

```

// ApplyPayload sends the payload to update_engine for
// application.
// The payload URL is passed as a file:// or content:// URI.
ApplyPayload(ctx context.Context, req ApplyPayloadRequest)
    error

// MonitorProgress watches update_engine status updates via
// callbacks.
MonitorProgress(ctx context.Context, callback
    ProgressCallback) error

// SetActiveSlot marks the newly updated slot as active for
// next boot.
SetActiveSlot(ctx context.Context, slot string) error

// Reboot triggers device reboot into the new slot.
Reboot(ctx context.Context) error
}

// ApplyPayloadRequest contains the parameters for update_engine.
type ApplyPayloadRequest struct {
    PayloadURL    string    // file:///data/ota/download.bin
    PayloadProps  []string // from payload_properties.txt
    Headers       []string // e.g., "USER_AGENT=HelixOTA/1.0.0"
    Offset        int64    // offset into the payload file
    Size          int64    // size of the payload
}

// ProgressCallback receives status updates from update_engine.
type ProgressCallback func(status UpdateEngineStatus)

// UpdateEngineStatus mirrors update_engine's status codes.
type UpdateEngineStatus struct {
    State int // 0=IDLE, 1=CHECKING,
        2=UPDATE_AVAILABLE, 3=DOWNLOADING,
        // 4=INSTALLING, 5=POSTINSTALL,
        6=FINALIZING
    Progress float64 // 0.0 - 1.0
    CurrentOp string // e.g., "VERIFY", "SOURCE_COPY",
        "REPLACE_BZ"
    ErrorMessage string // non-empty on error
    IsIntermittent bool // true if update can be interrupted
        and resumed
}

```

3.2.5 State Machine

The client update lifecycle is governed by a strict state machine that ensures no invalid transitions can occur.

```

// UpdateState represents the device's current update state.
type UpdateState string

const (

```

```

    StateIdle      UpdateState = "IDLE"
    StateChecking  UpdateState = "CHECKING"
    StateAvailable UpdateState = "AVAILABLE"
    StateDownloading UpdateState = "DOWNLOADING"
    StateVerifying UpdateState = "VERIFYING"
    StateInstalling UpdateState = "INSTALLING"
    StateRebooting  UpdateState = "REBOOTING"
    StateCommitting UpdateState = "COMMITTING"
    StateSuccess    UpdateState = "SUCCESS"
    StateFailed     UpdateState = "FAILED"
    StateRollback   UpdateState = "ROLLBACK"
)

// Valid state transitions (from → to)
var validTransitions = map[UpdateState][]UpdateState{
    StateIdle:      {StateChecking},
    StateChecking:  {StateAvailable, StateIdle, StateFailed},
    StateAvailable: {StateDownloading, StateIdle, StateFailed},
    StateDownloading: {StateVerifying, StateIdle, StateFailed},
    StateVerifying: {StateInstalling, StateFailed},
    StateInstalling: {StateRebooting, StateFailed},
    StateRebooting:  {StateCommitting, StateRollback},
    StateCommitting: {StateSuccess, StateRollback},
    StateSuccess:    {StateIdle},
    StateFailed:     {StateIdle},
    StateRollback:   {StateIdle, StateFailed},
}

// StateMachine enforces valid state transitions.
type StateMachine struct {
    current UpdateState
    mu       sync.Mutex
}

// Transition attempts to move to the next state.
// Returns ErrInvalidTransition if the transition is not allowed.
func (sm *StateMachine) Transition(next UpdateState) error {
    sm.mu.Lock()
    defer sm.mu.Unlock()


    allowed, ok := validTransitions[sm.current]
    if !ok {
        return fmt.Errorf("no transitions from state %s",
            sm.current)
    }

    for _, s := range allowed {
        if s == next {
            sm.current = next
            return nil
        }
    }
}

```

}

State Diagram (Mermaid-ready):



3.2.6 Report Service

Responsibility: Status reporting back to server with offline queuing and retry logic.

```
// ReportService sends device status events to the OTA server.
```

```
type ReportService struct {
    client    *http.Client
    serverURL string
    deviceID  string
    queue     []TelemetryEvent // offline queue, max 1000 events
    mu        sync.Mutex
}
```

```
// Report sends a status event to the server.  
// If the server is unreachable, the event is queued locally.
```

```
func (rs *ReportService) Report(
    ctx context.Context,
    eventType string,
    artifactID string,
    progress int,
    err error,
) error {
    event := TelemetryEvent{
        DeviceID:      rs.deviceID,
        ArtifactID:    artifactID,
        EventType:     eventType,
        ProgressPercent: progress,
        Timestamp:     time.Now(),
    }

    if err != nil {
        errMsg := err.Error()
        event.ErrorCode = ptrString(extractErrorCode(err))
        event.ErrorMessage = &errMsg
    }
}
```

```

    // Try to send immediately
    if sendErr := rs.send(ctx, event); sendErr != nil {
        rs.enqueue(event)
        return sendErr
    }

    // Try to flush queued events
    go rs.flushQueue(context.Background())

    return nil
}

func (rs *ReportService) enqueue(event TelemetryEvent) {
    rs.mu.Lock()
    defer rs.mu.Unlock()
    if len(rs.queue) >= 1000 {
        rs.queue = rs.queue[1:] // drop oldest
    }
    rs.queue = append(rs.queue, event)
}

```

3.3 Dashboard Components

3.3.1 Upload UI

Drag-and-drop OTA zip upload with real-time validation feedback. The upload flow uses chunked multipart upload with progress indicators:

- **Drag Zone:** Accepts .zip files only; validates file extension on drop
- **Metadata Form:** Pre-upload fields: target_version, target_model (dropdown), min_source_version
- **Upload Progress:** Chunked upload with byte-level progress bar; shows upload speed and ETA
- **Validation Feedback:** Real-time display of server-side validation stages (hash → signature → structure → compatibility); each stage shows a checkmark or error icon as it completes
- **Result View:** After successful upload: displays artifact ID, SHA-256 hash, file size, and validation status

3.3.2 Rollout Manager

Create, edit, pause, resume, and monitor rollouts:

- **Create Rollout Form:** Artifact selector (dropdown of valid artifacts), device group selector, initial percentage slider (1-100%), increment step selector, auto-advance toggle, mandatory flag, deadline date picker
- **Percentage Slider:** Visual slider with discrete stops at 1%, 5%, 10%, 25%, 30%, 50%, 75%, 100%; real-time preview of estimated device count at each percentage
- **Rollout Status Card:** Shows current/target percentage, success/failure/in-progress counts, progress bar, pause/resume button, advance button

- **Rollout Timeline:** Historical view of percentage advances with timestamps and health metrics at each stage

3.3.3 Device Fleet View

Table and card views showing all registered devices:

- **Table View:** Sortable columns: device ID, serial, model, current version, status, last check-in, group; filterable by status, group, version; paginated (50 per page)
- **Card View:** Grid of device cards showing key info at a glance; color-coded by status (green=up-to-date, blue=updating, red=failed, gray=offline)
- **Device Detail:** Click-through to device history: all telemetry events, update history, error details

3.3.4 Telemetry Dashboard

Charts showing update metrics over time:

- **Update Success Rate:** Line chart (time series) showing % of successful updates per rollout per day
- **Failure Reasons:** Bar chart showing top 10 error codes with device count
- **Rollout Progress:** Stacked area chart showing devices in each state (downloading, installing, success, failed) over time
- **Check-in Frequency:** Histogram showing device check-in intervals; highlights stale devices
- **Download Performance:** Distribution chart of download durations and sizes

3.3.5 Auth/Login UI

Secure authentication with TOTP 2FA:

- **Login Form:** Username + password; validates against server-side bcrypt hash
- **2FA Challenge:** After successful password, shows TOTP code input; supports recovery codes
- **Session Management:** JWT stored in httpOnly cookie; 30-minute inactivity timeout; refresh token rotation
- **Role-Based Navigation:** UI elements hidden/disabled based on user role (admin sees user management, operator sees rollout controls, viewer sees read-only dashboards)

4. Data Flow Diagrams

4.1 Update Check Flow

```
sequenceDiagram
    participant D as Android Device
    participant S as OTA Server
    participant C as Redis Cache
    participant DB as PostgreSQL
```

```

D->>S: GET /api/v1/devices/{id}/update-check
S->>C: GET device:{id}
alt Cache Hit
    C-->>S: Device data
else Cache Miss
    S->>DB: SELECT * FROM devices WHERE id = ?
    DB-->>S: Device data
    S->>C: SET device:{id} (TTL 60s)
end
S->>C: GET rollout:active:{group}
alt Cache Hit
    C-->>S: Rollout data
else Cache Miss
    S->>DB: SELECT * FROM rollouts WHERE group = ? AND status
= 'ACTIVE'
    DB-->>S: Rollout data
    S->>C: SET rollout:active:{group} (TTL 60s)
end
S->>S: Compute cohort: fnv32(device_id) % 100
alt Cohort < rollout_percentage
    S->>DB: SELECT * FROM artifacts WHERE id = ?
    DB-->>S: Artifact data
    S-->>D: 200 OK {update_available: true, ...}
else Cohort >= rollout_percentage
    S-->>D: 204 No Content
end

```

4.2 Upload and Validation Flow

sequenceDiagram

```

participant UI as Dashboard
participant S as OTA Server
participant W as Worker Pool
participant ST as MinIO/S3
participant DB as PostgreSQL

```

```

UI->>S: POST /api/v1/artifacts (multipart upload)
S->>S: Stream to temp file, compute SHA-256
S->>W: Submit validation task

```

```

W->>W: Stage 1: Hash verification
W->>W: Stage 2: RSA signature verification
W->>W: Stage 3: ZIP structure validation (payload.bin,
payload_properties.txt, care_map.pb)
W->>W: Stage 4: Compatibility check (model, version)

```

```

alt All stages pass
    W-->>S: ValidationResult{Valid: true}
    S->>S: Sign artifact with RSA-4096 private key
    S->>ST: Upload artifact to S3 bucket
    S->>DB: INSERT INTO artifacts (...)
    S->>S: Publish EventBus: artifact.uploaded
    S-->>UI: 201 Created {artifact_id, sha256,

```

```

validation_status: "VALID"}
    else Validation fails
        W-->>S: ValidationResult{Valid: false, Errors: [...]}
        S-->>DB: INSERT INTO artifacts (validation_status:
"INVALID")
        S-->>UI: 422 Unprocessable Entity {errors: [...]}
    end

```

4.3 Rollout Decision Flow

```

flowchart TD
    A[Rollout Created] --> B{Auto-advance?}
    B -->|No| C[Wait for manual advance]
    B -->|Yes| D[Start 15-min evaluation timer]
    D --> E{≥80% cohort reported?}
    E -->|No| D
    E -->|Yes| F{Failure rate < 5%?}
    F -->|No| G[Auto-pause rollout]
    G --> H[Send anomaly alert]
    F -->|Yes| I[Advance by increment_step]
    I --> J{Current ≥ 100%?}
    J -->|Yes| K[Mark rollout COMPLETED]
    J -->|No| D
    C --> L[Admin advances percentage]
    L --> F

```

4.4 Download and Install Flow

```

sequenceDiagram
    participant D as Android Device
    participant S as OTA Server
    participant ST as MinIO/S3 CDN
    participant UE as update_engine

    D->>S: GET /api/v1/devices/{id}/update-check
    S-->>D: 200 OK {download_url, sha256, signature_url}
    D->>S: GET /api/v1/artifacts/{id}/signature
    S-->>D: RSA-4096 signature bytes
    D->>D: State: AVAILABLE → DOWNLOADING
    D->>S: POST /api/v1/devices/{id}/status {status: DOWNLOADING}
    D->>ST: GET /artifacts/{id}/download (with Range header for
resume)
    ST-->>D: Artifact bytes (chunked)
    D->>D: State: DOWNLOADING → VERIFYING
    D->>S: POST /api/v1/devices/{id}/status {status: VERIFYING}
    D->>D: Verify SHA-256 hash
    D->>D: Verify RSA-4096 signature
    D->>D: Verify ZIP structure
    D->>D: State: VERIFYING → INSTALLING
    D->>S: POST /api/v1/devices/{id}/status {status: INSTALLING}
    D->>UE: ApplyPayload(payload_url, props, offset, size)
    UE-->>D: onStatusUpdate(INSTALLING, progress%)

```



```

D->>S: POST /api/v1/devices/{id}/status {status: INSTALLING,
progress: 67}
UE-->>D: onStatusUpdate(POSTINSTALL, 100%)
D->>D: State: INSTALLING → REBOOTING
D->>S: POST /api/v1/devices/{id}/status {status: REBOOTING}
D->>UE: SetActiveSlot("_b")
D->>D: Reboot device
Note over D: Device boots into new slot
D->>D: State: REBOOTING → COMMITTING
D->>S: POST /api/v1/devices/{id}/status {status: COMMITTING}
D->>D: State: COMMITTING → SUCCESS
D->>S: POST /api/v1/devices/{id}/status {status: SUCCESS,
version: "15.0.1-20260315"}

```

4.5 Telemetry Reporting Flow

flowchart TD

```

A[Device generates telemetry event] --> B{Server reachable?}
B -->|Yes| C[POST /api/v1/devices/{id}/status]
C --> D[Server ingests event]
D --> E[Publish to EventBus]
E --> F[Metrics Aggregator]
E --> G[Anomaly Detector]
E --> H[Notification Dispatcher]
F --> I[Update Redis cache: rollout health]
G --> J{Anomaly detected?}
J -->|Yes| K[Auto-pause rollout + alert]
J -->|No| L[Continue monitoring]
H --> M[Push WebSocket event to dashboard]
B -->|No| N[Enqueue locally - max 1000 events]
N --> O[Retry with backoff]
O --> B

```

5. Android 15 / RK3588 Specific Architecture

5.1 A/B Partition Layout on RK3588

The Orange Pi 5 Max with RK3588 uses the following A/B partition layout for Android 15:

RK3588 GPT Partition Table			
Partition	Slot A	Slot B	Shared
bootloader	bootloader_a	bootloader_b	—
boot	boot_a	boot_b	—
init_boot	init_boot_a	init_boot_b	—
vendor_boot	vendor_boot_a	vendor_boot_b	—
dtbo	dtbo_a	dtbo_b	—

vbmata	vbmata_a	vbmata_b	–
super	–	–	super (dynamic)
system	(dynamic_a)	(dynamic_b)	–
vendor	(dynamic_a)	(dynamic_b)	–
product	(dynamic_a)	(dynamic_b)	–
odm	(dynamic_a)	(dynamic_b)	–
misc	–	–	misc
userdata	–	–	userdata
metadata	–	–	metadata
cache	–	–	cache
fw_esp	–	–	fw_esp

Key RK3588 Considerations: - The super partition uses dynamic partitions (Android 10+). System, vendor, product, and odm are logical partitions within super, managed by libsnapshot and libdm. - init_boot is mandatory for Android 13+ and contains the generic ramdisk (previously bundled in boot). - vbmeta partitions contain Android Verified Boot metadata and must be updated atomically with their corresponding boot/system partitions. - The misc partition stores boot control information: active slot, retry count, and boot state. The boot control HAL reads and writes this partition. - RK3588's eMMC typically provides 32GB–256GB storage. With dynamic partitions and Virtual A/B, the COW snapshot space is carved from super partition's free space plus the userdata partition's snapshot logical volume.

5.2 update_engine Daemon Integration via Binder

The Android client integrates with update_engine through Binder IPC. The integration layer is implemented in Kotlin/Java and provides bindings to the Go SDK:

```
// UpdateEngineBridge.kt – Bridges Go SDK calls to Android
update_engine
class UpdateEngineBridge(
    private val context: Context,
    private val updateEngine: IUpdateEngine // Binder proxy
) : InstallOrchestrator {

    private val statusCallback = object :
        IUpdateEngineCallback.Stub() {
        override fun onStatusUpdate(status: Int, progress:
            Double) {
            // Map update_engine status codes to Go SDK
            UpdateEngineStatus
            val mapped = UpdateEngineStatus(
                state = mapEngineStatus(status),
                progress = progress,
            )
            // Forward to Go SDK via gomobile callback
            nativeOnStatusUpdate(mapped.toByteArray())
        }

        override fun onPayloadApplicationComplete(errorCode: Int)
        {
```

```

        if (errorCode ==
UpdateEngineErrorCodeConstants.SUCCESS) {
            nativeOnApplyComplete(true, "")
        } else {
            nativeOnApplyComplete(false, "update_engine
error: $errorCode")
        }
    }
}

override fun applyPayload(
    url: String,
    props: List<String>,
    offset: Long,
    size: Long
) {
    // Register callback before applying
    updateEngine.bind(statusCallback)

    // Apply payload via Binder IPC
    updateEngine.applyPayload(url, props.toTypedArray(),
        offset, size)
}

private fun mapEngineStatus(status: Int): Int {
    return when (status) {
        UpdateEngineStatusConstants.IDLE -> 0
        UpdateEngineStatusConstants.CHECKING_FOR_UPDATE -> 1
        UpdateEngineStatusConstants.UPDATE_AVAILABLE -> 2
        UpdateEngineStatusConstants.DOWNLOADING -> 3
        UpdateEngineStatusConstants.VERIFYING -> 4
        UpdateEngineStatusConstants.FINALIZING -> 5
        UpdateEngineStatusConstants.UPDATED_NEED_REBOOT -> 6
        else -> 0
    }
}
}
}

```

Binder IPC Flow: 1. Helix Android Client calls `IUpdateEngine.applyPayload()` via Binder proxy 2. `update_engine` daemon receives the request and begins streaming the payload 3. Status updates arrive via `IUpdateEngineCallback.onStatusUpdate()` Binder callback 4. Completion is signaled via `IUpdateEngineCallback.onPayloadApplicationComplete()` 5. The client then calls `IBootControl.setActiveBootSlot()` via another Binder interface

5.3 Boot Control HAL Interaction

The boot control HAL (`android.hardware.boot@1.2`) manages slot selection and boot state:

```

// BootControl interface abstracts the Android boot control HAL.
// On RK3588, this is implemented by Rockchip's boot control HAL.

```

```

type BootControl interface {
    // GetActiveSlot returns the currently active slot index (0 or
    // 1).
    GetActiveSlot() int

    // GetCurrentSlot returns the slot the device booted from.
    GetCurrentSlot() int

    // SetActiveBootSlot marks the given slot as active for next
    // boot.
    SetActiveBootSlot(slot int) error

    // SetSlotAsUnbootable marks a slot as unbootable (used on
    // failure).
    SetSlotAsUnbootable(slot int) error

    // IsSlotBootable returns whether the given slot is marked
    // bootable.
    IsSlotBootable(slot int) bool

    // GetSuffix returns the suffix for a slot ("_a" or "_b").
    GetSuffix(slot int) string
}

```

RK3588 Boot Control Implementation Notes: - Rockchip provides android.hardware.boot@1.2-impl-rockchip in their BSP - The HAL writes to the misc partition to persist slot selection across reboots - The bootloader (U-Boot on RK3588) reads misc partition to determine which slot to boot - If a slot fails to boot N times (typically 3), U-Boot falls back to the other slot - The retry count is decremented by the HAL on each successful boot and reset after commit

5.4 Virtual A/B (dm-snapshot) Considerations

Android 15 on RK3588 uses Virtual A/B, which introduces snapshot-based merging:

Virtual A/B Update Flow (RK3588)	
1. Pre-update:	super partition: [system_a][vendor_a][product_a][free] Device boots from slot A partitions.
2. During update:	COW snapshots created for each partition being updated. update_engine writes new data to COW, not to live data. super partition: [system_a][COW_sys][vendor_a][COW_vnd]
3. After reboot:	Device boots from COW snapshots (dm-snapshot). Old data is still intact in original partitions.
4. Merge phase (background):	snapshotctl merges COW data into original partitions.

Can be interrupted and resumed on next boot. After merge: super partition: [system_b][vendor_b][free]
5. Rollback (if boot fails): COW snapshots are discarded. Device boots from original (unchanged) slot A partitions.

Helix OTA Virtual A/B Integration: - The client does NOT directly manage dm-snapshot or COW — update_engine handles this internally - The client must ensure sufficient COW space is available before starting the update (check libsnapshot APIs) - After reboot, the client monitors the merge phase via snapshotctl and reports merge progress - If the merge is interrupted (e.g., power loss), the client reports MERGE_INTERRUPTED and the system resumes merge on next boot automatically

5.5 Rockchip-Specific Update Mechanisms

RK3588 has some Rockchip-specific considerations:

Concern	Standard Android	RK3588/Rockchip	Helix Adaptation
Bootloader update	Standard A/B bootloader	Rockchip upgrade_tool for bootloader/firmware	Bootloader updates handled separately via upgrade_tool; Helix manages system/ vendor/ product OTA only
Resource partition	Not applicable	resource_a/b for kernel DTB and logo	Included in OTA payload if resource partition changes
TrustZone firmware	Standard TEE	Rockchip OP-TEE in trust_a/b	OP-TEE updates must be coordinated with system updates
Rockchip rkboot	Not applicable	Rockchip-specific boot image format	boot_a/b uses Rockchip’s mkbootimg format
upgrade_tool	Not used	CLI tool for low-level flash operations	Used for initial provisioning and bootloader recovery; not for OTA updates

5.6 OTA Zip Payload Format

The Helix OTA artifact is a standard Android OTA zip with the following structure:

```
update.zip
├─ payload.bin                # Main update payload
(DeltaArchiveManifest + data blobs)
├─ Header (26 bytes)
└─ Magic: "CrAU" (4 bytes)
```

```

├── Version: uint64 (e.g., 2)
├── Manifest length: uint64
├── Metadata signature length: uint32
├── Padding
├── Manifest (protobuf)
│   ├── PartitionUpdate[] (per partition)
│   │   ├── partition_name: "system", "vendor", "boot", etc.
│   │   ├── old_partition_info: {sha256, size}
│   │   ├── new_partition_info: {sha256, size}
│   │   ├── InstallOperation[]
│   │   │   ├── type: REPLACE, REPLACE_BZ, REPLACE_XZ, ZERO,
│   │   │   │   SOURCE_COPY, SOURCE_BSDIFF, PUFFDIFF
│   │   │   ├── data_offset, data_length
│   │   │   └── src_extents[], dst_extents[]
│   │   └── max_timestamp, dynamic_partition_metadata
│   ├── Metadata signature (RSA-4096)
│   └── Data blobs (operation data)
├── payload_properties.txt # Update engine properties
│   ├── FILE_HASH=sha256_hex
│   ├── FILE_SIZE=bytes
│   ├── METADATA_HASH=sha256_hex
│   └── METADATA_SIZE=bytes
├── care_map.pb # Care map for dm-verity post-
install
│   └── CareMap { partitions: [{name, extents}] }
├── META-INF/
│   └── com/
│       ├── android/
│       │   └── otacert # RSA certificate for package
verification

```

5.7 Hash File Verification Integration

The Helix client verifies artifact integrity through a multi-layer approach:

```
// RK3588Verification performs RK3588-specific verification steps.
func RK3588Verification(artifactPath string, info *UpdateInfo)
    error {
    // Layer 1: Whole-file SHA-256 (computed during download)
    // Already verified by DownloadManager.

    // Layer 2: RSA-4096 signature on payload metadata
    // Verifies that the payload manifest has not been tampered
    with.
    // This is the same verification that update_engine performs.
    if err := verifyPayloadSignature(artifactPath,
        info.Signature); err != nil {
        return fmt.Errorf("payload signature: %w", err)
    }

    // Layer 3: payload_properties.txt hash verification
    // The properties file contains hashes of the payload data and
    metadata.
```

```

props, err := parsePayloadProperties(artifactPath)
if err != nil {
    return fmt.Errorf("parse properties: %w", err)
}
if err := verifyPayloadPropsHashes(artifactPath, props); err != nil {
    return fmt.Errorf("properties hash: %w", err)
}

// Layer 4: care_map.pb validation
// Ensures dm-verity care map is present and well-formed.
if err := validateCareMap(artifactPath); err != nil {
    return fmt.Errorf("care map: %w", err)
}

return nil
}

```

6. Integration with Existing Submodules

The Helix OTA server is built on the vasic-digital submodule ecosystem, adhering to the principle of never re-implementing solved problems. Each submodule is integrated at a specific layer of the architecture.

6.1 auth — JWT Token Management, RBAC

Integration Point: AuthService wraps vasic-digital/auth for all authentication and authorization operations.

```

// Integration with vasic-digital/auth
import "vasic-digital/auth"

type AuthService struct {
    jwtProvider auth.JWTProvider
    rbac         auth.RBACProvider
}

// The auth submodule provides:
// - JWTProvider: token issuance, validation, refresh, rotation
// - RBACProvider: role definitions, permission checks, role assignment
// - TokenBlacklist: revoked token tracking (Redis-backed)

// Usage pattern:
func (s *AuthService) IssueDeviceToken(ctx context.Context,
    deviceID string) (string, error) {
    claims := auth.Claims{
        Subject:  deviceID,
        Issuer:   "helix-ota",
        ExpiresAt: time.Now().Add(24 * time.Hour),
    }
}

```

```

        Custom: map[string]interface{}{
            "type":      "device",
            "device_id": deviceID,
        },
    }
    return s.jwtProvider.Issue(ctx, claims)
}

func (s *AuthService) CheckPermission(ctx context.Context, role
    Role, perm Permission) bool {
    return s.rbac.HasPermission(ctx, string(role), string(perm))
}

```

6.2 database — PostgreSQL Adapter, Migrations, Repository Pattern

Integration Point: All repository implementations use vasic-digital/database for connection pooling, migrations, and query execution.

```

// Integration with vasic-digital/database
import "vasic-digital/database"

// Database configuration
type DBConfig struct {
    Host      string    `env:"DB_HOST"
    default:"localhost"`
    Port      int       `env:"DB_PORT"
    default:"5432"`
    Name      string    `env:"DB_NAME"
    default:"helix_ota"`
    User      string    `env:"DB_USER"
    default:"helix"`
    Password  string    `env:"DB_PASSWORD"
    required:"true"`
    MaxOpenConns int     `env:"DB_MAX_OPEN"
    default:"25"`
    MaxIdleConns int     `env:"DB_MAX_IDLE"
    default:"10"`
    ConnMaxLifetime time.Duration `env:"DB_CONN_LIFETIME"
    default:"5m"`
}

// Migration files are managed by the database submodule
// and stored in: migrations/
// 001_create_devices.up.sql
// 001_create_devices.down.sql
// 002_create_artifacts.up.sql
// ...

// Repository pattern using database.Pool
type DeviceRepositoryImpl struct {
    db *database.Pool
}

```



```

func (r *DeviceRepositoryImpl) GetByID(ctx context.Context, id
    string) (*Device, error) {
    var device Device
    err := r.db.QueryRow(ctx,
        "SELECT id, serial, model, current_version, slot_suffix,
        hardware_fingerprint, device_group, status, last_check_in,
        created_at, updated_at FROM devices WHERE id = $1",
        id,
    ).Scan(&device.ID, &device.Serial, &device.Model,
        &device.CurrentVersion,
        &device.SlotSuffix, &device.HardwareFingerprint,
        &device.Group,
        &device.Status, &device.LastCheckIn, &device.CreatedAt,
        &device.UpdatedAt)
    if err != nil {
        return nil, fmt.Errorf("query device: %w", err)
    }
    return &device, nil
}

```

6.3 cache — Two-Level Caching for Device Inventory and Update Metadata

Integration Point: vasic-digital/cache provides a two-level cache (L1 in-process, L2 Redis) used for device inventory, rollout state, and update check results.

```

// Integration with vasic-digital/cache
import "vasic-digital/cache"

// Two-level cache configuration
// L1: In-process sync.Map with TTL eviction (sub-millisecond reads)
// L2: Redis with configurable TTL (millisecond reads, shared across server instances)
// Read path: L1 → L2 → Database → populate L2 → populate L1
// Write path: Invalidate L1 + L2 → Write Database

// Cache key design for device inventory
const (
    CacheKeyDevice      = "device:%s"           // L1 TTL: 30s,
    L2 TTL: 60s
    CacheKeyRollout     = "rollout:active:%s"    // L1 TTL: 15s,
    L2 TTL: 60s
    CacheKeyArtifact    = "artifact:%s"         // L1 TTL: 5m,
    L2 TTL: 15m
    CacheKeyHealth      = "rollout:health:%s"    // L1 TTL: 30s,
    L2 TTL: 60s
    CacheKeyUpdateInfo  = "update:%s:%s"        // L1 TTL: 30s,
    L2 TTL: 60s
)

// Usage in UpdateService

```

```

func (s *UpdateService) CheckForUpdate(ctx context.Context,
    deviceID string) (*UpdateInfo, error) {
    cacheKey := fmt.Sprintf(CacheKeyUpdateInfo, deviceID,
        "latest")

    // Try L1+L2 cache
    if cached, err := s.cache.Get(ctx, cacheKey); err == nil {
        return cached.(*UpdateInfo), nil
    }

    // Cache miss: compute result
    info, err := s.computeUpdateInfo(ctx, deviceID)
    if err != nil {
        return nil, err
    }
    if info != nil {
        s.cache.Set(ctx, cacheKey, info, 60*time.Second)
    }

    return info, nil
}

```

6.4 observability — Distributed Tracing, Metrics, Structured Logging

Integration Point: vasic-digital/observability provides OpenTelemetry tracing, Prometheus metrics, and structured JSON logging across all services.

```

// Integration with vasic-digital/observability
import "vasic-digital/observability"

// Tracing: every API handler starts an OpenTelemetry span
func (h *UpdateHandler) CheckForUpdate(w http.ResponseWriter, r
    *http.Request) {
    ctx, span := observability.StartSpan(r.Context(),
        "UpdateService.CheckForUpdate")
    defer span.End()

    deviceID := chi.URLParam(r, "deviceID")
    span.SetAttributes(attribute.String("device.id", deviceID))

    info, err := h.service.CheckForUpdate(ctx, deviceID)
    if err != nil {
        span.RecordError(err)
        // ...
    }
}

// Metrics: custom Prometheus metrics for OTA-specific operations
var (
    updateCheckDuration =
        observability.NewHistogram("helix_update_check_duration_seconds",
            "Duration of update check requests")

```

```

artifactUploadDuration =
    observability.NewHistogram("helix_artifact_upload_duration_seconds",
        "Duration of artifact upload + validation")
rolloutAdvanceTotal =
    observability.NewCounter("helix_rollout_advance_total",
        "Number of rollout percentage advances")
deviceStatusReported =
    observability.NewCounterVec("helix_device_status_reported_total",
        "Device status reports by status type", "status")
)

// Structured logging with request ID propagation
func (s *UpdateService) CheckForUpdate(ctx context.Context,
    deviceID string) (*UpdateInfo, error) {
    logger := observability.FromContext(ctx)
    logger.Info("checking for update",
        "device_id", deviceID,
        "trace_id", observability.TraceID(ctx),
    )
    // ...
}

```

6.5 security — TLS Configuration, Certificate Management

Integration Point: vasic-digital/security provides TLS 1.3 configuration, certificate loading, and mTLS verification for device-server communication.

```

// Integration with vasic-digital/security
import "vasic-digital/security"

// Server TLS configuration
func configureServerTLS(cfg *Config) (*tls.Config, error) {
    tlsConfig, err := security.NewTLSConfig(&security.TLSOptions{
        CertFile:      cfg.TLSCertPath,
        KeyFile:       cfg.TLSKeyPath,
        MinVersion:    tls.VersionTLS13,
        ClientAuth:    tls.VerifyClientCertIfGiven, // mTLS for
        // devices
        ClientCAFile:  cfg.DeviceCAPath,           // CA for
        // device certificates
        CurvePreferences: []tls.CurveID{
            tls.X25519,
            tls.CurveP256,
        },
    })
    if err != nil {
        return nil, fmt.Errorf("configure TLS: %w", err)
    }
    return tlsConfig, nil
}

// Device certificate verification

```

```

func verifyDeviceCertificate(cert *x509.Certificate)
    (*DeviceIdentity, error) {
    // CN format: helix-device-{device_id}
    if !strings.HasPrefix(cert.Subject.CommonName, "helix-
        device-") {
        return nil, security.ErrInvalidCertificate
    }
    deviceID := strings.TrimPrefix(cert.Subject.CommonName,
        "helix-device-")
    return &DeviceIdentity{DeviceID: deviceID, AuthMethod:
        "mTLS"}, nil
}

```

6.6 middleware — HTTP Middleware Chain

Integration Point: vasic-digital/middleware provides the HTTP middleware chain applied to all API routes.

```

// Integration with vasic-digital/middleware
import "vasic-digital/middleware"

func setupMiddleware(r chi.Router, cfg *Config) {
    r.Use(middleware.RequestID)           // X-Request-ID
        propagation
    r.Use(middleware.RealIP)              // X-Forwarded-For
        handling
    r.Use(middleware.Logger)              // Structured request
        logging
    r.Use(middleware.Recoverer)           // Panic recovery → 500
    r.Use(middleware.Timeout(30*time.Second)) // Request timeout
    r.Use(middleware.RateLimit(cfg.RateLimit)) // vasic-digital/
        ratelimiter
    r.Use(middleware.CORS(cfg.CORSAllowedOrigins))
    r.Use(middleware.Auth(authService)) // JWT/mTLS auth
        extraction
}

```

6.7 config — Configuration Loading

Integration Point: vasic-digital/config provides structured configuration loading from environment variables, config files, and command-line flags.

```

// Integration with vasic-digital/config
import "vasic-digital/config"

type Config struct {
    Server      ServerConfig  `env_prefix:"HELIX_SERVER_"`
    Database    DBConfig      `env_prefix:"HELIX_DB_"`
    Storage     StorageConfig `env_prefix:"HELIX_STORAGE_"`
    Auth        AuthConfig   `env_prefix:"HELIX_AUTH_"`
    Cache       CacheConfig  `env_prefix:"HELIX_CACHE_"`
}

```

```

type ServerConfig struct {
    Host      string    `env:"HOST"
               default:"0.0.0.0"`
    Port      int       `env:"PORT"
               default:"8443"`
    TLSCertPath string   `env:"TLS_CERT"
               required:"true"`
    TLSKeyPath string   `env:"TLS_KEY"
               required:"true"`
    DeviceCAPath string  `env:"DEVICE_CA"
               required:"true"`
    ReadTimeout time.Duration `env:"READ_TIMEOUT"
               default:"15s"`
    WriteTimeout time.Duration `env:"WRITE_TIMEOUT"
               default:"60s"`
}

// Load reads configuration from environment variables, falling
// back to
// helix_ota.yaml config file, falling back to defaults.
func Load() (*Config, error) {
    return config.Load[Config](
        config.WithEnvFile(".env"),
        config.WithConfigFile("helix_ota.yaml"),
        config.WithEnvPrefix("HELIX"),
    )
}

```

6.8 EventBus — Internal Event Publishing for Update Lifecycle

Integration Point: vasic-digital/EventBus provides publish/subscribe for decoupled, asynchronous event processing across services.

```

// Integration with vasic-digital/EventBus
import "vasic-digital/eventbus"

// Event types in the Helix OTA system
const (
    EventDeviceRegistered = "device.registered"
    EventDeviceStatusChanged = "device.status_changed"
    EventArtifactUploaded = "artifact.uploaded"
    EventRolloutCreated = "rollout.created"
    EventRolloutAdvanced = "rollout.advanced"
    EventRolloutPaused = "rollout.paused"
    EventAnomalyDetected = "anomaly.detected"
    EventTelemetryIngested = "telemetry.ingested"
)

// EventBus wiring in main()
func wireEventBus(bus eventbus.EventBus, services *Services) {
    // Telemetry events → metrics aggregation + anomaly detection
    bus.Subscribe("telemetry.*", services.Telemetry.ProcessEvent)
}

```

```

// Rollout events → notification dispatch → dashboard push
bus.Subscribe("rollout.*",
    services.Notification.HandleRolloutEvent)

// Artifact events → delta generation trigger (future: 1.0.2)
bus.Subscribe("artifact.uploaded",
    services.Artifact.HandleUploadEvent)

// Anomaly events → auto-pause rollout + alert
bus.Subscribe("anomaly.detected",
    services.Rollout.HandleAnomaly)
}

```

6.9 Storage — Artifact Storage (S3-Compatible)

Integration Point: vasic-digital/storage provides S3-compatible object storage for OTA artifacts, supporting multipart upload for large files.

```

// Integration with vasic-digital/storage
import "vasic-digital/storage"

// Storage configuration
type StorageConfig struct {
    Provider    string `env:"PROVIDER"      default:"minio" //
               "minio", "s3", "local"
    Endpoint    string `env:"ENDPOINT"      default:"localhost:9000"
    Bucket      string `env:"BUCKET"        default:"helix-ota-
               artifacts"
    AccessKey   string `env:"ACCESS_KEY"    required:"true"
    SecretKey   string `env:"SECRET_KEY"    required:"true"
    Region      string `env:"REGION"        default:"us-east-1"
    UseTLS      bool  `env:"USE_TLS"       default:"true"
}

// Usage in ArtifactService
func (s *ArtifactService) UploadToStorage(ctx context.Context,
    key string, path string) error {
    // Multipart upload for files > 64MB
    info, err := os.Stat(path)
    if err != nil {
        return err
    }
    if info.Size() > 64*1024*1024 {
        return s.storage.MultipartUpload(ctx, s.bucket, key, path,
            storage.WithPartSize(16*1024*1024),
            storage.WithConcurrency(4),
        )
    }
    return s.storage.Upload(ctx, s.bucket, key, path)
}

```

```

func (s *ArtifactService) GenerateDownloadURL(ctx
    context.Context, key string) (string, error) {
    // Generate pre-signed URL with 1-hour expiry for direct
    // download
    return s.storage.PresignedGet(ctx, s.bucket, key, time.Hour)
}

```

6.10 ratelimiter — API Rate Limiting

Integration Point: vasic-digital/ratelimiter provides per-device and per-API rate limiting.

```

// Integration with vasic-digital/ratelimiter
import "vasic-digital/ratelimiter"

// Rate limit configuration
var rateLimits = map[string]ratelimiter.Limit{
    // Device endpoints: 1 update check per minute per device
    "POST:/api/v1/devices/{id}/update-check": {Rate: 1, Burst: 2,
        Window: time.Minute},

    // Status reporting: 10 reports per minute per device
    "POST:/api/v1/devices/{id}/status": {Rate: 10, Burst: 15,
        Window: time.Minute},

    // Artifact upload: 5 per hour per user
    "POST:/api/v1/artifacts": {Rate: 5, Burst: 2, Window:
        time.Hour},

    // Rollout management: 30 per minute per user
    "POST:/api/v1/rollouts": {Rate: 30, Burst: 10, Window:
        time.Minute},

    // General API: 100 per minute per IP
    "*": {Rate: 100, Burst: 20, Window: time.Minute},
}

// Rate limiting middleware
func RateLimitMiddleware(limiter *ratelimiter.RateLimiter)
    func(http.Handler) http.Handler {
    return func(next http.Handler) http.Handler {
        return http.HandlerFunc(func(w http.ResponseWriter, r
            *http.Request) {
            key := rateLimitKey(r) // device_id, user_id, or IP
            limit := rateLimits[rateLimitRoute(r)]

            if !limiter.Allow(key, limit) {
                http.Error(w, "rate limit exceeded",
                    http.StatusTooManyRequests)
                return
            }
            next.ServeHTTP(w, r)
        })
    }
}

```

```

    }
}

```

6.11 concurrency — Worker Pools for Parallel Artifact Processing

Integration Point: vasic-digital/concurrency provides goroutine pool management for CPU-intensive operations like artifact validation.

```

// Integration with vasic-digital/concurrency
import "vasic-digital/concurrency"

// Worker pool configuration
func NewArtifactWorkerPool() *concurrency.Pool {
    return concurrency.NewPool(concurrency.PoolConfig{
        MaxWorkers:    runtime.NumCPU(),    // CPU-bound: match
        // CPU count
        MaxQueueSize:  100,                  // Buffer up to 100
        // validation tasks
        IdleTimeout:   30 * time.Second,    // Kill idle workers
        // after 30s
        PanicHandler: func(v interface{}) { // Graceful panic
            // recovery
            observability.Global().Error("worker panic",
                "value", v,
                "stack", string(debug.Stack()),
            )
        },
    })
}

// Usage: parallel SHA-256 computation during upload
func (s *ArtifactService) streamToTempWithHash(
    ctx context.Context,
    reader io.Reader,
) (string, string, int64, error) {
    tempFile, err := os.CreateTemp("", "helix-upload-*)
    if err != nil {
        return "", "", 0, err
    }

    hash := sha256.New()
    size, err := io.Copy(tempFile, io.TeeReader(reader, hash))
    if err != nil {
        os.Remove(tempFile.Name())
        return "", "", 0, err
    }

    sha256Hex := hex.EncodeToString(hash.Sum(nil))
    return tempFile.Name(), sha256Hex, size, nil
}

```


6.12 recovery — Circuit Breakers for External Service Calls

Integration Point: `vasic-digital/recovery` provides circuit breakers for calls to external services (S3, Redis, database) and graceful shutdown handling.

```
// Integration with vasic-digital/recovery
import "vasic-digital/recovery"

// Circuit breakers for external dependencies
func wireCircuitBreakers() map[string]*recovery.CircuitBreaker {
    return map[string]*recovery.CircuitBreaker{
        "storage": recovery.NewCircuitBreaker(recovery.Config{
            Name:          "artifact-storage",
            MaxFailures: 5,           // Open after 5 consecutive
                             failures
            Timeout:      30 * time.Second, // Try again after 30s
            HalfOpenMax: 1,           // Allow 1 test request in
                             half-open
        }),
        "cache": recovery.NewCircuitBreaker(recovery.Config{
            Name:          "redis-cache",
            MaxFailures: 10,
            Timeout:      10 * time.Second,
            HalfOpenMax: 3,
        }),
    }
}

// Usage in ArtifactService
func (s *ArtifactService) UploadToStorage(ctx context.Context,
    key, path string) error {
    cb := s.circuitBreakers["storage"]
    return cb.Execute(func() error {
        return s.storage.Upload(ctx, s.bucket, key, path)
    })
}

// If S3 is down, the circuit breaker opens, and UploadToStorage
// returns
// recovery.ErrCircuitOpen immediately without attempting the
// call.
// The artifact upload endpoint returns 503 Service Unavailable.

// Graceful shutdown
func (s *Server) Shutdown(ctx context.Context) error {
    // recovery.GracefulShutdown handles:
    // 1. Stop accepting new connections
    // 2. Drain in-flight requests (30s timeout)
    // 3. Close database connections
    // 4. Flush remaining EventBus events
    // 5. Close WebSocket connections with close frame
    return recovery.GracefulShutdown(ctx, s.httpServer, s.db,
        s.eventBus)
}
```

7. Communication Protocols

7.1 REST API (Primary Client-Server Protocol)

The REST API is the primary communication channel between devices and the server. It follows OpenAPI 3.1 specification with JSON request/response bodies.

Endpoint	Method	Purpose	Auth	Rate Limit
/api/v1/devices	POST	Device registration	mTLS/ JWT	5/min
/api/v1/devices/{id}/update-check	GET	Check for available update	mTLS/ JWT	1/min
/api/v1/devices/{id}/status	POST	Report device status	mTLS/ JWT	10/min
/api/v1/artifacts	POST	Upload OTA artifact	JWT (admin/ operator)	5/hour
/api/v1/artifacts/{id}/download	GET	Download artifact	mTLS/ JWT	10/min
/api/v1/artifacts/{id}/signature	GET	Download artifact signature	mTLS/ JWT	10/min
/api/v1/rollouts	POST	Create rollout	JWT (admin/ operator)	30/min
/api/v1/rollouts/{id}	GET	Get rollout status	JWT	60/min
/api/v1/rollouts/{id}	PATCH	Update rollout (% , pause/ resume)	JWT (admin/ operator)	30/min
/api/v1/telemetry/rollout/{id}	GET	Get rollout health metrics	JWT	60/min
/api/v1/auth/login	POST	Dashboard login	None	10/min
/api/v1/auth/2fa	POST	TOTP 2FA verification	Session	5/min

Design Principles: - All endpoints use HTTPS (TLS 1.3 mandatory) - Device endpoints use mTLS when available, falling back to JWT - Dashboard endpoints use JWT with short-lived access tokens (15 min) and long-lived refresh tokens (7 days) - Idempotency: POST /status is idempotent by (device_id, artifact_id, timestamp) composite key - Pagination: List endpoints support ?offset=0&limit=50 with Link headers - Error format: RFC 7807 Problem Details (application/problem+json)

7.2 WebSocket/SSE (Real-Time Dashboard Updates)

The dashboard receives real-time updates via WebSocket (primary) with SSE fallback:

```
// WebSocket protocol (JSON messages)
// Client → Server:
{
  "type": "subscribe",
  "filters": ["device_group:rk3588_opi5max",
             "rollout:rol_01HXYZ"]
}

// Server → Client:
{
  "type": "device.status_changed",
  "payload": {
    "device_id": "dev_01HABC",
    "status": "INSTALLING",
    "progress": 67,
    "timestamp": "2026-03-05T12:34:56Z"
  }
}

{
  "type": "rollout.advanced",
  "payload": {
    "rollout_id": "rol_01HXYZ",
    "old_percentage": 10,
    "new_percentage": 30,
    "health": {"failure_rate": 0.02, "success_count": 150}
  }
}
```

7.3 gRPC (Internal Microservice Communication)

While the MVP is a monolith, internal service boundaries use gRPC interfaces for future decomposition:

```
// proto/helix/ota/v1/telemetry.proto
service TelemetryService {
  rpc IngestTelemetry(IngestTelemetryRequest) returns
    (IngestTelemetryResponse);
  rpc GetRolloutHealth(GetRolloutHealthRequest) returns
    (RolloutHealth);
}
```

```

    rpc StreamDeviceEvents(StreamDeviceEventsRequest) returns
        (stream DeviceEvent);
}

message IngestTelemetryRequest {
    repeated TelemetryEvent events = 1;
}

message TelemetryEvent {
    string device_id = 1;
    string artifact_id = 2;
    string event_type = 3;
    string status = 4;
    int32 progress_percent = 5;
    optional string error_code = 6;
    optional string error_message = 7;
    google.protobuf.Timestamp timestamp = 8;
}

```

7.4 Binder IPC (Android update_engine Communication)

On the Android device, the Helix client communicates with update_engine through Binder IPC. This is handled by the native Android layer:

Interface	Direction	Purpose
IUpdateEngine.applyPayload()	Client → Engine	Initiate update application
IUpdateEngine.bind()	Client → Engine	Register status callback
IUpdateEngine.suspend()	Client → Engine	Pause update
IUpdateEngine.resume()	Client → Engine	Resume paused update
IUpdateEngine.cancel()	Client → Engine	Cancel in-progress update
IUpdateEngineCallback.onStatusUpdate()	Engine → Client	Progress notifications
IUpdateEngineCallback.onPayloadApplicationComplete()	Engine → Client	Completion/failure notification
IBootControl.setActiveBootSlot()	Client → HAL	Mark new slot as active
IBootControl.getCurrentSlot()	Client → HAL	Query current boot slot

8. Security Architecture

8.1 Security Overview

The Helix OTA system implements defense-in-depth across all communication layers and data stores. Security is not a feature bolted on after the fact — it is an architectural constraint that shapes every design decision.

Security Layers	
Layer 5: Artifact Integrity	• RSA-4096 signature on every OTA payload • SHA-256 hash verification at download and install time • Server-side validation chain rejects unsigned artifacts
Layer 4: Payload Authenticity	• Signing key stored in HSM (production) / vault (dev) • Signing occurs in isolated CI pipeline; key never on OTA • Client embeds trusted public key; rejects unknown keys
Layer 3: Transport Security	• TLS 1.3 mandatory on all connections • Certificate pinning on device client • HSTS headers on all HTTP responses
Layer 2: Device Authentication	• mTLS with per-device X.509 certificates • JWT fallback with 24-hour rotation • Nonce-based request signing prevents replay
Layer 1: Application Authorization	• RBAC with admin/operator/viewer roles • Per-endpoint permission checks • TOTP 2FA for dashboard access • Rate limiting per device and per user

8.2 TLS 1.3 for All Communications

Connection	TLS Version	Certificate	Notes
Device → Server	TLS 1.3	Server cert + device client cert (mTLS)	Device cert CN = helix-device-{id}
Dashboard → Server	TLS 1.3	Server cert	JWT in httpOnly cookie

Connection	TLS Version	Certificate	Notes
Server → PostgreSQL	TLS 1.3	Internal CA cert	Database requires client cert
Server → Redis	TLS 1.3	Internal CA cert	Redis 7+ native TLS
Server → MinIO/S3	TLS 1.3	S3 endpoint cert	AWS or MinIO CA
Server → Server (gRPC)	TLS 1.3	Internal CA cert	mTLS between service instances

8.3 RSA-4096 Signing of OTA Artifacts

```
// Artifact signing pipeline (runs in isolated CI environment)
func SignArtifact(artifactPath string, privateKey
    *rsa.PrivateKey) ([]byte, error) {
    // 1. Hash the payload.bin file
    payload, err := extractPayloadBin(artifactPath)
    if err != nil {
        return nil, err
    }
    hash := sha256.Sum256(payload)


    // 2. Sign the hash with RSA-4096
    signature, err := rsa.SignPKCS1v15(
        rand.Reader,
        privateKey,
        crypto.SHA256,
        hash[:],
    )
    if err != nil {
        return nil, fmt.Errorf("sign payload: %w", err)
    }

    return signature, nil
}
```

Signing Key Management: - **Development:** RSA-4096 key stored in HashiCorp Vault - **Production:** RSA-4096 key stored in HSM (CloudHSM or on-premises Nitrokey HSM) - Key rotation: New key pair generated annually; old public key remains embedded in client for backward compatibility - Key ceremony: Multi-party authorization required for production key access

8.4 mTLS for Device-Server Authentication





Device certificates are issued during initial provisioning (factory or first boot). The device generates a CSR locally, sends it to the server over an initial secure channel, and receives a signed certificate valid for 1 year. Certificate renewal happens automatically 30 days before expiry.

8.5 JWT for Dashboard API Authentication

```

// JWT token structure
type DashboardClaims struct {
    UserID    string `json:"uid"`
    Username  string `json:"usr"`
    Role      Role   `json:"role"`
    jwt.RegisteredClaims
}

// Access token: 15-minute expiry, httpOnly cookie
// Refresh token: 7-day expiry, httpOnly cookie, single-use
//              (rotated on refresh)
// Token rotation: refresh token is invalidated after use; new
//              pair issued

func (s *AuthService) Login(ctx context.Context, username,
    password, totpCode string) (*TokenPair, error) {
    // 1. Verify password (bcrypt)
    user, err := s.userRepo.GetByUsername(ctx, username)
    if err != nil {
        return nil, ErrInvalidCredentials
    }
    if err :=
        bcrypt.CompareHashAndPassword([]byte(user.PasswordHash),
            []byte(password)); err != nil {
        return nil, ErrInvalidCredentials
    }

    // 2. Verify TOTP code
    if !totp.Validate(totpCode, user.TOTPSecret) {
        return nil, ErrInvalid2FA
    }

    // 3. Issue token pair
    accessToken, err := s.jwtProvider.Issue(ctx, DashboardClaims{
        UserID:    user.ID,
        Username:  user.Username,
        Role:      user.Role,
  
```

```

    RegisteredClaims: jwt.RegisteredClaims{
        ExpiresAt: jwt.NewNumericDate(time.Now().Add(15 *
            time.Minute)),
        IssuedAt:   jwt.NewNumericDate(time.Now()),
        Issuer:     "helix-ota",
    },
})

refreshToken, err := s.jwtProvider.IssueRefresh(ctx, user.ID)

return &TokenPair{
    AccessToken:  accessToken,
    RefreshToken: refreshToken,
}, nil
}

```

8.6 RBAC with Three Roles

Permission	Admin	Operator	Viewer
Upload artifact	✓	✓	✗
Delete artifact	✓	✗	✗
Create rollout	✓	✓	✗
Update/pause rollout	✓	✓	✗
Delete rollout	✓	✗	✗
View devices	✓	✓	✓
Delete device	✓	✗	✗
View telemetry	✓	✓	✓
Manage users	✓	✗	✗
View dashboard	✓	✓	✓

8.7 Artifact Upload Validation Chain

Upload Request



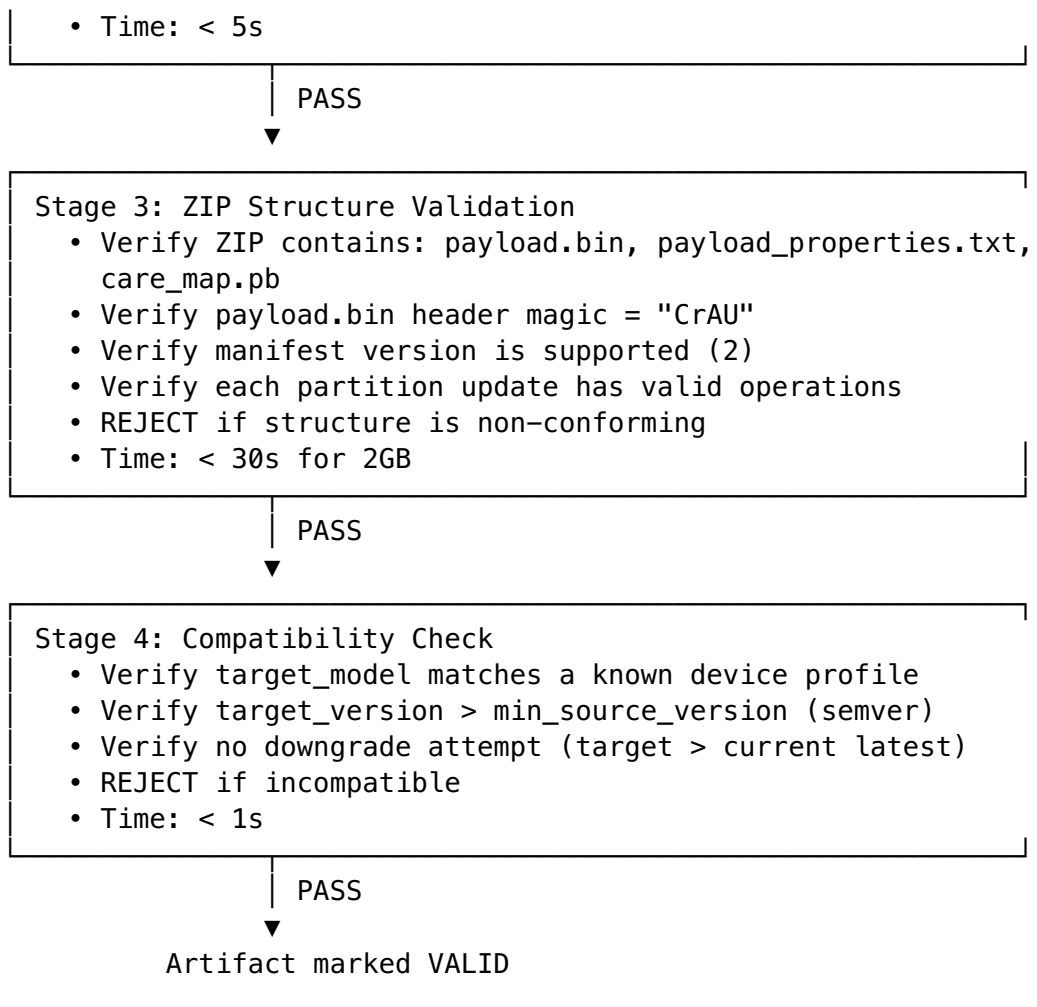
Stage 1: SHA-256 Hash Verification

- Compute SHA-256 of uploaded file
- Compare with hash reported by client
- REJECT if mismatch (corruption or tampering)
- Time: $O(n)$ in file size, ~30s for 2GB

PASS

Stage 2: RSA Signature Verification


- Verify payload.bin signature against trusted public key
- REJECT if signature invalid or missing




9. Scalability Design

9.1 Horizontal Scaling of OTA Server

The OTA server is designed for horizontal scaling behind a load balancer:






Scaling Characteristics:

- **Stateless servers:** All server instances are stateless. Device session state is stored in Redis, not in server memory. Any server can handle any device request.
- **Shared-nothing architecture:** No server-to-server communication required for request handling. Communication happens only through shared data stores (PostgreSQL, Redis) and EventBus (Redis pub/sub).
- **Auto-scaling:** Kubernetes Horizontal Pod Autoscaler (HPA) scales based on CPU (target: 60%) and request latency (target: p95 < 500ms).
- **Typical capacity:** A single Go server instance handles ~5,000 concurrent update check requests at p95 < 200ms. For a fleet of 100,000 devices checking every 4 hours, 2-3 server instances suffice.

9.2 CDN-Based Artifact Distribution

For large fleets, artifact downloads are offloaded to a CDN:



CDN Strategy:

- Artifacts are immutable: once uploaded, they never change. Perfect for CDN caching.
- Cache key: `/ {artifact_id} / {filename}` — guaranteed unique per artifact version
- Cache TTL: 30 days (artifacts are never invalidated)
- Pre-signed URLs: Server generates time-limited (1 hour) pre-signed download URLs for CDN, preventing unauthorized access
- Range request support: CDN must support HTTP Range headers for resumable downloads
- Estimated bandwidth savings: 95%+ cache hit rate for large fleets (many devices downloading the same artifact)

9.3 Database Sharding Strategy

For MVP, a single PostgreSQL instance with read replicas is sufficient. For scale beyond 500,000 devices, the following sharding strategy applies:

Table	Sharding Strategy	Rationale
devices	Hash-shard by id (e.g., 16 shards)	Uniform distribution; most queries are by device ID
telemetry_events	Range-partition by timestamp (monthly)	

Table	Sharding Strategy	Rationale
artifacts	Single shard (small table, <10K rows)	Time-series data; old partitions can be archived/dropped Artifacts are few and frequently accessed
rollouts	Single shard (small table, <1K rows)	Rollouts are few and rarely written
device_status	Hash-shard by device_id	Co-located with devices shard for JOIN efficiency

Read Replicas: - Analytical queries (telemetry aggregation, dashboard metrics) route to read replicas - Write queries (device check-in, status report) route to primary - Replication lag target: < 100ms (acceptable for update checks; stale data is corrected on next check-in)

9.4 WebSocket Connection Management at Scale

The notification service must handle thousands of concurrent dashboard WebSocket connections:

```
// WebSocketHub manages connections with sharding and graceful
// cleanup.
type WebSocketHub struct {
    shards    []*wsShard // Shard connections by user ID
    hash
    shardCount int // Number of shards (default:
    runtime.NumCPU())
}

type wsShard struct {
    mu sync.RWMutex
    sessions map[string]*DashboardSession // session_id →
    session
    register chan *DashboardSession
    unregister chan string // session_id
    broadcast chan *NotificationEvent
}

// Shard assignment: session is assigned to shard = fnv32(user_id)
// % shardCount
// This ensures all sessions for the same user are on the same
// shard,
// enabling efficient targeted broadcasting.

// Health monitoring: stale connections (no ping for 60s) are
// evicted
func (s *wsShard) evictStale() {
    s.mu.Lock()
    defer s.mu.Unlock()
    for id, session := range s.sessions {
```

```

        if time.Since(session.LastPing) > 60*time.Second {
            session.Conn.Close()
            delete(s.sessions, id)
        }
    }
}

```

WebSocket Scaling at a Glance:

Metric	Single Instance	3-Instance Cluster
Concurrent connections	~10,000	~30,000
Messages per second	~50,000	~150,000
Memory per connection	~8 KB	~8 KB
Total memory (10K connections)	~80 MB	~80 MB per instance

Cross-Instance Broadcasting: When the EventBus publishes a notification event, it is broadcast via Redis Pub/Sub to all server instances. Each instance's WebSocketHub receives the event and delivers it to locally-connected sessions matching the event's filter criteria.

Appendix A: Database Schema (Key Tables)

-- Devices table: core device registry

```

CREATE TABLE devices (
    id          VARCHAR(32) PRIMARY KEY,
    serial      VARCHAR(64) NOT NULL UNIQUE,
    model       VARCHAR(64) NOT NULL,
    current_version VARCHAR(64) NOT NULL,
    target_version VARCHAR(64),
    slot_suffix VARCHAR(4) NOT NULL DEFAULT '_a',
    hardware_fingerprint VARCHAR(128) NOT NULL UNIQUE,
    device_group VARCHAR(64) NOT NULL DEFAULT 'default',
    status      VARCHAR(32) NOT NULL DEFAULT 'IDLE',
    last_check_in TIMESTAMPTZ NOT NULL,
    created_at   TIMESTAMPTZ NOT NULL DEFAULT NOW(),
    updated_at   TIMESTAMPTZ NOT NULL DEFAULT NOW()
);

```

```

CREATE INDEX idx_devices_group ON devices(device_group);
CREATE INDEX idx_devices_status ON devices(status);
CREATE INDEX idx_devices_last_check_in ON devices(last_check_in);

```

-- Artifacts table: OTA update packages

```

CREATE TABLE artifacts (
    id          VARCHAR(32) PRIMARY KEY,
    filename     VARCHAR(256) NOT NULL,
    target_version VARCHAR(64) NOT NULL,
    target_model  VARCHAR(64) NOT NULL,
    min_source_version VARCHAR(64) NOT NULL,

```

```

sha256          VARCHAR(64) NOT NULL,
size_bytes      BIGINT NOT NULL,
storage_key     VARCHAR(512) NOT NULL,
signature       BYTEA NOT NULL,
validation_status VARCHAR(16) NOT NULL DEFAULT 'PENDING',
created_at      TIMESTAMPTZ NOT NULL DEFAULT NOW()
);

CREATE INDEX idx_artifacts_model ON artifacts(target_model);
CREATE INDEX idx_artifacts_status ON artifacts(validation_status);

-- Rollouts table: phased deployment records
CREATE TABLE rollouts (
    id          VARCHAR(32) PRIMARY KEY,
    artifact_id VARCHAR(32) NOT NULL REFERENCES
        artifacts(id),
    device_group VARCHAR(64) NOT NULL,
    current_percentage INTEGER NOT NULL DEFAULT 0,
    target_percentage INTEGER NOT NULL DEFAULT 100,
    increment_step  INTEGER NOT NULL DEFAULT 10,
    auto_advance    BOOLEAN NOT NULL DEFAULT FALSE,
    status          VARCHAR(16) NOT NULL DEFAULT 'ACTIVE',
    mandatory       BOOLEAN NOT NULL DEFAULT FALSE,
    deadline        TIMESTAMPTZ,
    created_at      TIMESTAMPTZ NOT NULL DEFAULT NOW(),
    updated_at      TIMESTAMPTZ NOT NULL DEFAULT NOW()
);

CREATE INDEX idx_rollouts_status ON rollouts(status);
CREATE INDEX idx_rollouts_group ON rollouts(device_group);

-- Telemetry events table: device status reports
CREATE TABLE telemetry_events (
    id          VARCHAR(32) PRIMARY KEY,
    device_id   VARCHAR(32) NOT NULL REFERENCES
        devices(id),
    artifact_id VARCHAR(32) REFERENCES artifacts(id),
    event_type  VARCHAR(32) NOT NULL,
    status      VARCHAR(32) NOT NULL,
    progress_percent INTEGER NOT NULL DEFAULT 0,
    error_code  VARCHAR(32),
    error_message TEXT,
    metadata    JSONB,
    timestamp   TIMESTAMPTZ NOT NULL,
    ingested_at TIMESTAMPTZ NOT NULL DEFAULT NOW()
);

-- Partitioned by month for efficient time-range queries and
  archival
CREATE INDEX idx_telemetry_device ON telemetry_events(device_id);
CREATE INDEX idx_telemetry_artifact ON
    telemetry_events(artifact_id);

```

```
CREATE INDEX idx_telemetry_type ON telemetry_events(event_type);
CREATE INDEX idx_telemetry_timestamp ON
    telemetry_events(timestamp);
```

Appendix B: Key Go Interfaces Summary

```
// Core repository interfaces (implemented via vasic-digital/
    database)
type DeviceRepository interface {
    GetByID(ctx context.Context, id string) (*Device, error)
    GetByFingerprint(ctx context.Context, fp string) (*Device,
        error)
    Create(ctx context.Context, device *Device) error
    Update(ctx context.Context, device *Device) error
    ListByGroup(ctx context.Context, group string, offset, limit
        int) ([]*Device, error)
    CountByGroup(ctx context.Context, group string) (int, error)
}

type ArtifactRepository interface {
    GetByID(ctx context.Context, id string) (*Artifact, error)
    Create(ctx context.Context, artifact *Artifact) error
    ListValid(ctx context.Context, model string) ([]*Artifact,
        error)
}

type RolloutRepository interface {
    GetByID(ctx context.Context, id string) (*Rollout, error)
    Create(ctx context.Context, rollout *Rollout) error
    Update(ctx context.Context, rollout *Rollout) error
    FindActiveForGroup(ctx context.Context, group string)
        (*Rollout, error)
}

type TelemetryRepository interface {
    BatchCreate(ctx context.Context, events []TelemetryEvent)
        error
    AggregateRolloutMetrics(ctx context.Context, rolloutID
        string) (*RolloutMetrics, error)
}

// Storage interface (implemented via vasic-digital/storage)
type StorageProvider interface {
    Upload(ctx context.Context, bucket, key string, reader
        io.Reader) error
    Download(ctx context.Context, bucket, key string)
        (io.ReadCloser, error)
    PresignedGet(ctx context.Context, bucket, key string, expiry
        time.Duration) (string, error)
    Delete(ctx context.Context, bucket, key string) error
}

// Cache interface (implemented via vasic-digital/cache)
```

```
type CacheProvider interface {  
    Get(ctx context.Context, key string) (interface{}, error)  
    Set(ctx context.Context, key string, value interface{}, ttl  
        time.Duration) error  
    Delete(ctx context.Context, key string) error  
}
```

Document Control: This document is a living specification. Any architectural change must be reflected here within 24 hours of implementation. Changes require review per HelixConstitution §3.2.