

sync_md_siblings

sync_md_siblings.sh — companion guide

Revision: 1 Last modified: 2026-06-11T00:00:00Z

Overview

scripts/sync_md_siblings.sh walks the §11.4.65 (“Universal Markdown export mandate”) INCLUDED Markdown corpus and (re)generates synchronized .html + .pdf siblings for every in-scope .md source. It is the out-of-the-box mechanism that keeps non-source-tree Markdown documents in sync across .md / .html / .pdf per §11.4.65 + §11.4.12.

The conversion pipeline is the sanctioned one documented in docs/qa/20260608-md-exports/README.md, applied per file:

```
pandoc <f>.md -s --metadata title=<base> -o <f>.html  
weasyprint <f>.html <f>.pdf
```

Prerequisites

- pandoc (tested with 3.9.0.2) — Markdown → standalone HTML.
- weasyprint (tested with 66.0) — HTML → PDF.
- timeout / gtimeout (optional) — bounds each conversion to 60s.

If pandoc or weasyprint is absent, the script logs an honest SKIP for the whole run and exits 0 **without faking any artifact** (§11.4.6 no-guessing). It never emits a fabricated or empty PDF.

Usage examples

```
# Generate missing/stale siblings across the included corpus:  
bash scripts/sync_md_siblings.sh
```

```
# Preview the work without writing anything:  
bash scripts/sync_md_siblings.sh --dry-run
```

```
# Force regeneration even when siblings are newer than sources:
bash scripts/sync_md_siblings.sh --force
```

The script resolves the repo root itself (`git rev-parse --show-toplevel`), so it can be invoked from any working directory.

Scope (what it includes / excludes)

INCLUDE (per §11.4.65):

- project-root *.md
- docs/**/*.*md — EXCLUDING per-run evidence dirs docs/qa/<run-id>/** (those evidence READMEs carry no siblings by established convention; `git ls-files 'docs/qa/**/README.html' | wc -l` is 0)
- scripts/**/*.*md
- owned-submodule top-level README.md / CLAUDE.md / AGENTS.md / CHANGELOG.md (constitution, containers, the six ota-* bricks, helixqa, challenges)

EXCLUDE: external/, prebuilts/, packages/modules/, kernel*/, out/, build/, node_modules/, .git/, third-party submodule internals (e.g. submodules/http3), and the docs/qa/<run-id>/ evidence dirs.

Source files are taken from git's tracked set (`git ls-files`), so untracked / ignored noise is never processed.

Internal behaviour

1. Argument parse (`--force`, `--dry-run`, `--help`).
2. Repo-root resolution + `cd`.
3. Dependency preflight — honest SKIP if pandoc/weasyprint missing.
4. Build the in-scope file list (4 include classes), de-duplicated.
5. Per file: skip when both siblings exist and are $\geq$ source mtime (unless `--force`); otherwise run pandoc then weasyprint.
6. Anti-bluff post-check — a file counts as generated ONLY if both .html and .pdf exist and are **non-empty** afterwards (§11.4.6).
7. Print SUMMARY: generated=N skipped=M failed=K; exit non-zero iff any failure occurred.

Each conversion is bounded by `timeout 60` (when available). Temp list files are removed on every exit path via `trap` (§11.4.14).

Edge cases

- **Missing tools** → whole-run SKIP, exit 0, nothing written.
- **weasyprint CSS warnings** on the pandoc default template are benign and do not fail the conversion (stderr is suppressed).
- **Newer siblings** are left untouched unless `--force`.
- **Empty post-conversion sibling** is reported FAIL, never counted as generated.

Related scripts

- docs/qa/20260608-md-exports/README.md — origin of the sanctioned pandoc+weasyprint pipeline.
- Project doc-sync wrappers (scripts/testing/sync_*) — the broader §11.4.12 / §11.4.53 / §11.4.59 export machinery this complements.

Last verified

2026-06-11 — run produced generated=52 skipped=99 failed=0; two generated PDFs spot-verified as valid (file → PDF document, version 1.7, non-zero size). pandoc 3.9.0.2 + weasyprint 66.0.